
MPI with Python

Gregor von Laszewski, laszewski@gmail.com

Fidel Leal

Jacques Fleischer

May 19, 2022

Contents

1	Preface	5
1.1	Acknowledgement	5
1.2	Online Availability	5
1.3	Document Management in GitHub	5
1.4	Document Notation	6
2	Introduction	6
3	Installation	7
3.1	Python Version	7
3.2	Operating Systems and MPI Versions	8
3.3	Getting the CPU Count	8
3.4	Windows 10 Home, Education, or Pro	8
3.5	macOS	10
3.6	Ubuntu	10
3.7	Raspberry Pi	11
3.8	Testing the Installation	11
4	Hosts, Machinefile, Rankfile	11
4.1	Running MPI on a Single Computer	11
4.2	Running MPI on Multiple Computers	12
4.2.1	Prerequisite	12
4.2.2	Using Hosts	12
4.2.3	Machinefile	13
4.2.4	Rankfiles for Multiple Cores	13
5	MPI Functionality	14
5.1	Differences to the C Implementation of MPI	14
5.1.1	Initialization	14
5.1.2	Capitalization for Pickle vs. Memory Messages	14
5.1.3	Using NumPy with mpi4py	15
5.2	MPI Functionality	15
5.2.1	Communicator	15
5.2.2	Point-to-Point Communication	15
5.3	Collective Communication	19
5.3.1	Broadcast	19
5.3.2	Scatter	22
5.3.3	Gather	24

5.3.4	Allgather Memory Objects	27
5.4	Process Management	28
5.4.1	Dynamic Process Management with spawn	28
5.4.2	Futures	30
6	MPI Example Programs	34
6.1	MPI Ring Example	34
6.2	Counting Numbers	36
6.3	Monte Carlo Calculation of Pi	37
6.3.1	Numba	41
6.3.2	Running Monte Carlo on multiple hosts	45
6.4	Mandelbrot	46
6.4.1	Assignments	48
6.5	Other MPI Example Programs	49
6.6	GPU Programming with MPI	49
7	Parameter Management	49
7.1	Using the Shell Variables to Pass Parameters	49
7.1.1	Using click to pass parameters	51
8	SLURM	52
8.1	Installation of SLURM on a Raspberry Pi Cluster	52
8.1.1	Method 1 - Install from Host	53
8.1.2	Method 2 - Install on Manager	53
8.2	Install SLURM on a Single Raspberry Pi	55
8.3	MPI Example	55
9	Links to Other Documents	56
9.1	Assignment	56
10	Appendix	56
10.1	Git Bash on Windows	56
10.2	Make on Windows	57
10.2.1	Installation	58
10.3	Installing WSL on Windows 10	58
10.4	Benchmarks	59
10.4.1	Prerequisites	60
10.4.2	System Parameters	60
10.4.3	Combining the logs	65
11	Assignments	65

References

67

1 Preface

Gregor von Laszewski has initiated this project as a voluntary summer research project available for university students.

1.1 Acknowledgement

Besides the coauthors, students who contributed to this document's very early version are Cooper Young, Erin Seliger, and Agness Lungua.

1.2 Online Availability

This document is in part published at:

- Medium <https://laszewski.medium.com/python-and-mpi-part-1-7e76a6ec1c6d>
- Friends Link: <https://laszewski.medium.com/python-and-mpi-part-1-7e76a6ec1c6d?sk=cc21262764659c0ef2d3ddc684f54034>

Please check them out as they may include slight improvements.

1.3 Document Management in GitHub

Note: The source document is managed at <https://cloudmesh.github.io/cloudmesh-mpi/doc/chapters>. To make changes or corrections please use a pull request.

The repository, documentation, and examples are available at:

- Repository: <https://github.com/cloudmesh/cloudmesh-mpi>
- Examples: <https://github.com/cloudmesh/cloudmesh-mpi/tree/main/examples>
- Documents:
 - <https://cloudmesh.github.io/cloudmesh-mpi/report-mpi.pdf>
 - <https://cloudmesh.github.io/cloudmesh-mpi/report-group.pdf>

To check out the repository use

```
1 $ git clone git@github.com:cloudmesh/cloudmesh-mpi.git
```

or

```
1 $ git clone https://github.com/cloudmesh/cloudmesh-mpi.git
```

In case you have `make` and `docker` installed on your machine, you can create this document locally with

```
1 $ make image
2 $ make
```

Please not that `make` can also be installed on Windows as documented in our appenix, so you can also create this document easily on Windows.

1.4 Document Notation

To keep things uniform, we use the following document notations.

1. Empty lines are to be placed before and after a context change, such as a headline, paragraph, list, image inclusion.
2. All code is written in code blocks using and three backquotes. A rendered example looks as follows:

```
1 a = "this is an example"
```

3. Single quote inclusion must be used for filenames and other names as they are referred to in code blocks.
4. To showcase command inclusion, we use a block but precede every command with a `$` or other prefix indicating the computer on which the command is executed.

```
1 $ ls
```

5. The Bibliography is for now managed via markdown footnotes or direct links

2 Introduction

Today Python [1] has become the predominantly programming language to coordinate scientific applications, especially machine and deep learning applications. However, previously existing parallel programming paradigms such as **Message Passing Interface (MPI)** have proven to be a useful asset when it comes to enhancing complex data flows while executing them on multiple computers , including supercomputers. The framework is well known in the C-language community. However, many practitioners do not have the time to learn C to utilize such advanced cyberinfrastructure. Hence, it is advantageous to access MPI from Python. We showcase how you can easily deploy and use MPI from Python via a tool called `mpi4py`.

Message Passing Interface (MPI) is a message-passing standard that allows for efficient data communication between the address spaces of multiple processes. The MPI standard began in 1992 as a collective effort by several organizations, institutions, vendors, and users. Since the first draft of the

specification in November 1993, the standard has undergone several revisions and updates leading to its current version: MPI 4.0 (June 2021).

Multiple implementations following the standard exist, including the two most popular MPICH [2] and OpenMPI [3]. However, other free or commercial implementations exist [[4]][[5]][[6]].

Additionally, MPI is a language-independent interface. Although support for C and Fortran is included as part of the standard, multiple libraries providing bindings for other languages are available, including those for Java, Julia, R, Ruby, and Python.

Thanks to its user-focused abstractions, its standardization, portability, and scalability, and availability MPI is a popular tool in the creation of high-performance and parallel computing programs.

3 Installation

Next, we discuss how to install mpi4py on various systems. We will focus on installing it on a single computer using multiple cores.

This Installation section does not cover the installation of SLURM, which is covered in a later section.

3.1 Python Version

In most cases you can probably use the newest version of Python and then add MPI for Python. However, we have currently only tested it for Python version 3.10.4. If you have tested it on newer versions, please let us know so we add it here to our compatibility list. While we have not tested it, we do not anticipate any issues running mpi4py on Windows 11.

Python Version	OS	Tested	Processor
3.10.4	Windows 10	yes	AMD
3.10.4	Windows 10	yes	Intel
3.9	Windows 10	yes	Intel
3.9	Mac	yes	Intel
3.9.7	Ubuntu 20.04	yes	AMD
3.9.7	Ubuntu 20.04	yes	Intel
3.10.4	RaspberryOS 11	yes	ARM
3.9.2	RaspberryOS 11	yes	ARM

3.2 Operating Systems and MPI Versions

The following table shows which operating systems use which version of MPI:

Operating System	MPI Version
Windows	MS-MPI v10.1.2
macOS	Open MPI v4.1.1
Ubuntu	MPICH v3.3.2
Raspberry Pi	Open MPI v4.1.0

3.3 Getting the CPU Count

For the examples listed in this document, knowing the number of cores on your computer is important. This can be found out through the command line or a python program.

In Python, you can do it with

```
1 import multiprocessing
2 multiprocessing.cpu_count()
```

or as a command line

```
1 $ python -c "import multiprocessing; print(multiprocessing.cpu_count())"
```

However, you can also use the command line tools that we have included in our documentation.

3.4 Windows 10 Home, Education, or Pro

1. We assume you have installed Git Bash on your computer. The installation is easy, but be careful to watch the various options at install time. Make sure it is added to the Path variable.

For details see: <https://git-scm.com/downloads>

2. We also assume you have installed Python3.9 according to either the installation at [python.org](https://www.python.org) or conda. We do recommend the installation from [python.org](https://www.python.org).

<https://www.python.org/downloads/>

You will need to install a python virtual env to avoid conflict by accident with your system installed version of Python.

For details on how to do this, please visit our extensive documentation at <https://cybertraining-dsc.github.io/docs/tutorial/reu/python/> under the subsection titled “Python venv”

3. Microsoft has its own implementation of MPI which we recommend at this time. First, you need to download msmpi from

- <https://docs.microsoft.com/en-us/message-passing-interface/microsoft-mpi#ms-mpi-downloads>

Go to the download link underneath the heading **MS-MPI Downloads** and download and install it. Select the two packages and click Next. When downloaded, click on them and complete the setups.

```
1 msmpisetup.exe
2 msmpisdk.msi
```

4. Open the system control panel and click on **Advanced system settings** (which can be searched for with the search box in the top-right, and then click **View advanced system settings**) and then click **Environment Variables...**

5. Under the user variables box, click on **Path**

6. Click New in order to add

C:\Program Files (x86)\Microsoft SDKs\MPI

and

C:\Program Files\Microsoft MPI\Bin

to the Path. The **Browse Directory...** button makes this easier, and the **Variable name** can correspond to each directory, e.g., “MPI” and “MPI Bin” respectively

7. Close any open bash windows and then open a new one

8. Type the command

```
1 $ which mpiexec
```

to verify if it works.

9. After you verified it is available, install mpi4py with

```
1 $ pip install mpi4py
```

ideally, while bash is in venv

10. Next, find out how many processes you can run on your machine and remember that number. You can do this with

```
1 $ wmic CPU Get DeviceID,NumberOfCores,NumberOfLogicalProcessors
```

Alternatively, you can use a python program as discussed in the section “Getting the CPU Count”

3.5 macOS

1. Find out how many processes you can run on your machine and remember that number. You can do this with

```
1 $ sysctl hw.physicalcpu hw.logicalcpu
```

2. First, install python 3 from <https://www.python.org/downloads/>
3. Next, install homebrew and install the open-mpi version of MPI as well as mpi4py:

```
1 $ xcode-select --install
2 $ /bin/bash -c "$(curl -fsSL https://raw.githubusercontent.com/
  Homebrew/install/HEAD/install.sh)"
3 $ brew install wget
4 $ brew install open-mpi
5 $ python3 -m venv ~/ENV3
6 $ source ~/ENV3/bin/activate
7 $ pip install mpi4py
```

If you are prompted to install command line developer tools, install them.

3.6 Ubuntu

These instructions apply to 20.04 and 21.04. Please use 20.04 in case you like to use GPUs.

1. First, find out how many processes you can run on your machine and remember that number. You can do this with

```
1 $ nproc
```

2. The installation of mpi4py on Ubuntu is relatively easy. Please follow these steps. We recommend that you create a python `venv` so you do not by accident interfere with your system python. As usual, you can activate it in your `.bashrc` file while adding the source line there. Lastly, make sure you check it out and adjust the `-n` parameters to the number of cores of your machine. In our example, we have chosen the number 4, you may have to change that value

```
1 $ sudo apt-get update
2 $ sudo apt install python3.9 python3.9-dev python3-dev python3.9-
  venv python3.8-venv
3 $ python3 -m venv ~/ENV3
4 $ source ENV3/bin/activate
5 (ENV3) $ sudo apt-get install -y mpich-doc mpich
6 (ENV3) $ pip install mpi4py -U
```

Any errors along the lines of

- `Python.h: No such file or directory` or
- `Could not build wheels for mpi4py which use PEP 517`

should be fixed by installing `python3-dev` in the `venv`

3.7 Raspberry Pi

1. Install Open MPI in your pi by entering the following command assuming a PI4, PI3B+ PI3, PI2:

```
1 $ python -m venv ~/ENV3
2 $ source ~/ENV3/bin/activate
3 $ sudo apt-get install openmpi-bin
4 $ mpicc --showme:version
5 $ pip install mpi4py
```

If you have other Raspberry Pi's you may need to update the core count according to the hardware specification.

3.8 Testing the Installation

On all systems, the installation is very easy. Just change in our example the number 4 to the number of cores in your system.

```
1 (ENV3) $ mpiexec -n 4 python -m mpi4py.bench helloworld
```

You will see an output similar to

```
1 Hello, World! I am process 0 of 4 on myhost.
2 Hello, World! I am process 1 of 4 on myhost.
3 Hello, World! I am process 2 of 4 on myhost.
4 Hello, World! I am process 3 of 4 on myhost.
```

where `myhost` is the name of your computer.

Note: the messages can be in a different order.

4 Hosts, Machinefile, Rankfile

4.1 Running MPI on a Single Computer

In case you like to try out MPI and just use it on a single computer with multiple cors, you can skip this section for now and revisit it, once you scale up and use multiple computers.

4.2 Running MPI on Multiple Computers

MPI is designed for running programs on multiple computers. One of these computers serves as manager and communicates to its workers. To define on which computer is running what, we need to have a configuration file that lists a number of hosts to participate in our set of machines, the MPI cluster.

The configuration file specifying this is called a machinefile or rankfile. We will explain the differences to them in this section.

4.2.1 Prerequisite

Naturally, the requisite to use a cluster is that you

1. have MPI and mpi4py installed on each of the computers, and
2. have access via ssh on each of these computers

If you use a Raspberry PI cluster, we recommend using our cloudmesh-pi-burn program [TODOREF]. This will conveniently create you a Raspberry PI cluster with login features established. You still need to install mpi4py, however on each node.

If you use another set of resources, you will often see the recommendation to use passwordless ssh key between the nodes. This we only recommend if you are an expert and have placed the cluster behind a firewall. If you experiment instead with your own cluster, we recommend that you use password-protected SSH keys on your manager node and populate them with ssh-copy-id to the worker computers. To not always have to type in your password to the different machines, we recommend you use `ssh-agent`, and `ssh-add`.

4.2.2 Using Hosts

In the case of multiple computers, you can simply specify the hosts as a parameter to your MPI program that you run on your manager node

```
1 (ENV3) $ mpiexec -n 4 -host re0,red1,red2,red3 python -m mpi4py.bench  
    helloworld
```

To specify how many processes you like to run on each of them, you can use the option `-ppn` followed by the number.

```
1 (ENV3) $ mpiexec -n 4 -ppn 2 -host re0,red1,red2,red3 python -m mpi4py.  
    bench helloworld
```

As today we usually have multiple cores on a processor, you could be using that core count as the parameter.

4.2.3 Machinefile

To simplify the parameter passing to MPI you can use machine files instead. This allows you also to define different numbers of processes for different hosts. Thus it is more flexible. In fact, we recommend that you use a machine file in most cases as you then also have a record of how you configured your cluster.

The machine file is a simple text file that lists all the different computers participating in your cluster. As MPI was originally designed at a time when there was only one core on a computer, the simplest machine file just lists the different computers. When starting a program with the machine file as option, only one core of the computer is utilized.

The machinefile can be explicitly passed along as a parameter while placing it in the manager machine

```
1 mpirun.openmpi \  
2 -np 2 \  
3 -machinefile /home/pi/mpi_testing/machinefile \  
4 python helloworld.py
```

An example of a simple machinefile contains the IP addresses. The username can be proceeded by the IP address.

```
1 pi@192.168.0.10:1  
2 pi@192.168.0.11:2  
3 pi@192.168.0.12:2  
4 pi@192.168.0.13:2  
5 pi@192.168.0.14:2
```

In many cases, your machine name may be available within your network and known to all hosts in the cluster. In that case, it is more convenient. To sue the machine names.

```
1 pi@red0:1  
2 pi@red1:2  
3 pi@red2:2  
4 pi@red3:2  
5 pi@red4:2
```

Please make sure to change the IP addresses or name of your hosts according to your network.

4.2.4 Rankfiles for Multiple Cores

In contrast to the host parameter, you can fine-tune the placement of processes to computers with a [rankfile](#). This may be important if your hardware has, for example specific computers for data storage or GPUs.

If you like to add multiple cores from a machine, you can also use a [rankfile](#)

```
1 mpirun -r my_rankfile --report-bindings ...
2
3 Where the rankfile contains:
4 rank 0=pi@192.168.0.10 slot=1:0
5 rank 1=pi@192.168.0.10 slot=1:1
6 rank 2=pi@192.168.0.11 slot=1:0
7 rank 3=pi@192.168.0.10 slot=1:1
```

In this configuration, we only use 2 cores from two different PIs.

5 MPI Functionality

In this section, we will discuss several useful MPI communication features.

5.1 Differences to the C Implementation of MPI

Before we start with a detailed introduction, we like to make those that have experience with non Python versions of MPI aware of some differences.

5.1.1 Initialization

In mpi4py, the standard `MPI_INIT()` and `MPI_FINALIZE()` commonly used to initialize and terminate the MPI environment are automatically handled after importing the mpi4py module. Although not generally advised, mpi4py still provides `MPI.Init()` and `MPI.Finalize()` for users interested in manually controlling these operations. Additionally, the automatic initialization and termination can be deactivated. For more information on this topic, please check the original mpi4py documentation:

- `MPI.Init()` and `MPI.Finalize()`
- Deactivating automatic initialization and termination on mpi4py

5.1.2 Capitalization for Pickle vs. Memory Messages

Another characteristic feature of mpi4py is the availability of uppercase and lowercase communication methods. Lowercase methods like `comm.send()` use Python's `pickle` module to transmit objects in a serialized manner. In contrast, the uppercase versions of methods like `comm.Send()` enable transmission of data contained in a contiguous memory buffer, as featured in the MPI standard. For additional information on the topic, the manual section Communicating Python Objects and Array Data.

5.1.3 Using NumPy with mpi4py

Several of the examples presented in the following sections use NumPy arrays to illustrate the behavior of mpi4py's uppercase communication methods.

NumPy is a Python library geared towards scientific computing. It features high-level mathematical functions that add support to work with and operate on multi-dimensional arrays and matrices.

NumPy quickly gained popularity thanks to its performance advantages in comparison to Python lists. NumPy array elements must have a uniform type and are stored contiguously in memory. As a consequence, memory consumption is lower and runtime performance improves, since there is no need to store type pointers or perform type checks before operating on any element. Type uniformity and contiguous memory use also allow for fast and efficient application of diverse mathematical operations to all indices of an array, making NumPy very attractive for use in statistical analysis, visualization libraries, and large data manipulation.

An interesting and useful exception to the type uniformity rule can be achieved by defining a NumPy array of Python objects, which allows for an array containing elements of different sizes/types, including other NumPy arrays.

To learn more about NumPy installation and use, please check our tutorials in Section 10.2 of Python for Cloud Computing.

5.2 MPI Functionality

5.2.1 Communicator

All MPI processes need to be addressable and are grouped in a `communicator`. The default communicator is called `world` and assigns a rank to each process within the communicator.

Thus all MPI programs we will discuss here start with

```
1 comm = MPI.COMM_WORLD
```

In the MPI program, the function

```
1 rank = comm.Get_rank()
```

returns the rank. This is useful to be able to write conditional programs that depend on the rank. Rank 0 is the rank of the manager process.

5.2.2 Point-to-Point Communication

5.2.2.1 Send and Receive Python Objects The `send()` and `recv()` methods provide for functionality to transmit data between two specific processes in the communicator group. It can be applied to

any Python data object that can be pickled. The advantage is that the object is preserved, however it comes with the disadvantage that pickling the data takes more time than a direct memory copy.



Figure 1: Sending and receiving data between two processes

Here is the definition for the `send()` method:

```
1 comm.send(buf, dest, tag)
```

`buf` represents the data to be transmitted, `dest` and `tag` are integer values that specify the rank of the destination process, and a tag to identify the message being passed, respectively. `tag` is particularly useful for cases when a process sends multiple kinds of messages to another process.

On the other end is the `recv()` method, with the following definition:

```
1 comm.recv(buf, source, tag, status)
```

In this case, `buf` can specify the location for the received data to be stored. In more recent versions of MPI, 'buf' has been deprecated. In those cases, we can simply assign `comm.recv(source, tag, status)` as the value of our buffer variable in the receiving process. Additionally, `source` and `tag` can specify the desired source and tag of the data to be received. They can also be set to `MPI.ANY_SOURCE` and `MPI.ANY_TAG`, or be left unspecified.

In the following example, an integer is transmitted from process 0 to process 1.

```
1 #!/usr/bin/env python
2 from mpi4py import MPI
3
4 # Communicator
5 comm = MPI.COMM_WORLD
6
7 # Get the rank of the current process in the communicator group
8 rank = comm.Get_rank()
9
10 # Variable to receive the data
11 data = None
12
13 # Process with rank 0 sends data to process with rank 1
14 if rank == 0:
15     comm.send(42, dest=1)
16
17 # Process with rank 1 receives and stores data
18 if rank == 1:
19     data = comm.recv(source=0)
```



```

20
21 # Each process in the communicator group prints its data
22 print(f'After send/receive, the value in process {rank} is {data}')
```

Executing `mpiexec -n 4 python send_receive.py` yields:

```

1 After send/receive, the value in process 2 is None
2 After send/receive, the value in process 3 is None
3 After send/receive, the value in process 0 is None
4 After send/receive, the value in process 1 is 42
```

As we can see, the transmission only occurred between processes 0 and 1, and no other process was affected.

5.2.2.2 Send and Recive Python Memory Objects The following example illustrates the use of the uppercase versions of the methods `comm.Send()` and `comm.Recv()` to perform a transmission of data between processes from memory to memory. In our example we will again be sending a message between processors of rank 0 and 1 in the communicator group.

```

1 #!/usr/bin/env python
2 import numpy as np
3 from mpi4py import MPI
4
5 # Communicator
6 comm = MPI.COMM_WORLD
7
8 # Get the rank of the current process in the communicator group
9 rank = comm.Get_rank()
10
11 # Create empty buffer to receive data
12 buf = np.zeros(5, dtype=int)
13
14 # Process with rank 0 sends data to process with rank 1
15 if rank == 0:
16     data = np.arange(1, 6)
17     comm.Send([data, MPI.INT], dest=1)
18
19 # Process with rank 1 receives and stores data
20 if rank == 1:
21     comm.Recv([buf, MPI.INT], source=0)
22
23 # Each process in the communicator group prints the content of its
    buffer
24 print(f'After Send/Receive, the value in process {rank} is {buf}')
```

Executing `mpiexec -n 4 python send_receive_buffer.py` yields:

```

1 After Send/Receive, the value in process 3 is [0 0 0 0 0]
2 After Send/Receive, the value in process 2 is [0 0 0 0 0]
3 After Send/Receive, the value in process 0 is [0 0 0 0 0]
```

```
4 After Send/Receive, the value in process 1 is [1 2 3 4 5]
```

5.2.2.3 Non-blocking Send and Receive MPI can also use non-blocking communications. This allows the program to send the message without waiting for the completion of the submission. This is useful for many parallel programs so we can overlap communication and computation while both take place simultaneously. The same can be done with receive, but if a message is not available and you do need the message, you may have to probe or even use a blocked receive. To wait for a message to be sent or received, we can also use the wait method, effectively converting the non-blocking message to a blocking one.

Next, we showcase an example of the non-blocking send and receive methods `comm.isend()` and `comm.irecv()`. Non-blocking versions of these methods allow for the processes involved in transmission/reception of data to perform other operations in overlap with the communication. In contrast, the blocking versions of these methods previously exemplified do not allow data buffers involved in transmission or reception of data to be accessed until any ongoing communication involving the particular processes has been finalized.

```
1 #!/usr/bin/env python
2 from mpi4py import MPI
3
4 # Communicator
5 comm = MPI.COMM_WORLD
6
7 # Get the rank of the current process in the communicator group
8 rank = comm.Get_rank()
9
10 # Variable to receive the data
11 data = None
12
13 # Process with rank 0 sends data to process with rank 1
14 if rank == 0:
15     send = comm.isend(42, dest=1)
16     send.wait()
17
18 # Process with rank 1 receives and stores data
19 if rank == 1:
20     receive = comm.irecv(source=0)
21     data = receive.wait()
22
23 # Each process in the communicator group prints its data
24 print(f'After isend/ireceive, the value in process {rank} is {data}')
```

Executing `mpirun -n 4 python isend_ireceive.py` yields:

```
1 After isend/ireceive, the value in process 2 is None
2 After isend/ireceive, the value in process 3 is None
3 After isend/ireceive, the value in process 0 is None
4 After isend/ireceive, the value in process 1 is 42
```

5.2.2.4 Ping Pong This example program uses the aforementioned `send()` and `recv()` methods to print a variable, `sendmsg`, depending on which rank the MPI program is presently working with.

```
1 from mpi4py import MPI
2
3 comm = MPI.COMM_WORLD
4 assert comm.size == 2
5
6 if comm.rank == 0:
7     sendmsg = 777
8     comm.send(sendmsg, dest=1, tag=55)
9     recvmsg = comm.recv(source=1, tag=77)
10 else:
11     recvmsg = comm.recv(source=0, tag=55)
12     sendmsg = "abc"
13     comm.send(sendmsg, dest=0, tag=77)
14 print(sendmsg)
```

This program can only be executed using `mpiexec -n 2 python pingpong.py`, which yields

```
1 abc
2 777
```

Note how, at first line-by-line glance, the program's code sets `sendmsg` to 777 before it is set to `abc`. However, upon program execution, the output is `abc` first because of the `dest` and `tag` values. On rank 0 (during program's initial stages), 777 is sent to destination 1. On rank 1 (remember there are only two ranks: 0 and 1), `abc` is sent to destination 0. The destination integers correspond to the ranks and the program leaves printing the `sendmsg` for last (after the `send()` and `recv()` methods have determined the variable values). This explains the output.

5.3 Collective Communication

5.3.1 Broadcast

The `bcast()` method and its memory version `Bcast()` broadcast a message from a specified *root* process to all other processes in the communicator group.

5.3.1.1 Broadcast of a Python Object In terms of syntax, `bcast()` takes the object to be broadcast and the parameter `root`, which establishes the rank number of the process broadcasting the data. If no `root` parameter is specified, `bcast` will default to broadcasting from the process with rank 0.

Thus, the two lines are functionally equivalent.

```
1 data = comm.bcast(data, root=0)
2 data = comm.bcast(data)
```

In our following example, we broadcast a two-entry Python dictionary from a root process to the rest of the processes in the communicator group.



Figure 2: Broadcasting data from a root process to the rest of the processes in the communicator group

The following code snippet shows the creation of the dictionary in process with rank 0. Notice how the variable `data` remains empty in all the other processes.

```
1 #!/usr/bin/env python
2 from mpi4py import MPI
3
4 # Set up the MPI Communicator
5 comm = MPI.COMM_WORLD
6
7 # Get the rank of the current process in the communicator group
8 rank = comm.Get_rank()
9
10 if rank == 0: # Process with rank 0 gets the data to be broadcast
11     data = {'size': [1, 3, 8],
12            'name': ['disk1', 'disk2', 'disk3']}
13 else: # Other processes' data is empty
14     data = None
15
16 # Print data in each process
17 print(f'before broadcast, data on rank {rank} is: {data}')
18
19 # Data from process with rank 0 is broadcast to other processes in our
20 # communicator group
21 data = comm.bcast(data, root=0)
22
23 # Print data in each process after broadcast
24 print(f'after broadcast, data on rank {rank} is: {data}')
```

After running `mpirun -n 4 python broadcast.py` we get the following:

```
1 before broadcast, data on rank 3 is: None
2 before broadcast, data on rank 0 is:
3   {'size': [1, 3, 8], 'name': ['disk1', 'disk2', 'disk3']}
4 before broadcast, data on rank 1 is: None
5 before broadcast, data on rank 2 is: None
6 after broadcast, data on rank 3 is:
```

```

7  {'size': [1, 3, 8], 'name': ['disk1', 'disk2', 'disk3']}
8  after broadcast, data on rank 0 is:
9  {'size': [1, 3, 8], 'name': ['disk1', 'disk2', 'disk3']}
10 after broadcast, data on rank 1 is:
11 {'size': [1, 3, 8], 'name': ['disk1', 'disk2', 'disk3']}
12 after broadcast, data on rank 2 is:
13 {'size': [1, 3, 8], 'name': ['disk1', 'disk2', 'disk3']}

```

As we can see, all other processes received the data broadcast from the root process.

5.3.1.2 Broadcast of a Memory Object In our following example, we broadcast a NumPy array from process 0 to the rest of the processes in the communicator group using the uppercase `comm.Bcast()` method.

```

1  #!/usr/bin/env python
2  import numpy as np
3  from mpi4py import MPI
4
5  # Communicator
6  comm = MPI.COMM_WORLD
7
8  # Get the rank of the current process in the communicator group
9  rank = comm.Get_rank()
10
11 # Rank 0 gets a NumPy array containing values from 0 to 9
12 if rank == 0:
13     data = np.arange(0, 10, 1, dtype='i')
14
15 # Rest of the processes get an empty buffer
16 else:
17     data = np.zeros(10, dtype='i')
18
19 # Print data in each process before broadcast
20 print(f'before broadcasting, data for rank {rank} is: {data}')
21
22 # Broadcast occurs
23 comm.Bcast(data, root=0)
24
25 # Print data in each process after broadcast
26 print(f'after broadcasting, data for rank {rank} is: {data}')

```

Executing `mpirun -n 4 python npbcast.py` yields:

```

1  before broadcasting, data for rank 1 is: [0 0 0 0 0 0 0 0 0 0]
2  before broadcasting, data for rank 2 is: [0 0 0 0 0 0 0 0 0 0]
3  before broadcasting, data for rank 3 is: [0 0 0 0 0 0 0 0 0 0]
4  before broadcasting, data for rank 0 is: [0 1 2 3 4 5 6 7 8 9]
5  after broadcasting, data for rank 0 is: [0 1 2 3 4 5 6 7 8 9]
6  after broadcasting, data for rank 2 is: [0 1 2 3 4 5 6 7 8 9]
7  after broadcasting, data for rank 3 is: [0 1 2 3 4 5 6 7 8 9]
8  after broadcasting, data for rank 1 is: [0 1 2 3 4 5 6 7 8 9]

```

As we can see, the values in the array at the process with rank 0 have been broadcast to the rest of the processes in the communicator group.

5.3.2 Scatter

While broadcast send all data to all processes, scatter send chunks of data to each process.

In our next example, we will `scatter` the members of a list among the processes in the communicator group. We illustrate the concept in the next figure, where we indicate the data that is scattered to the ranks processes with D_i




Figure 3: Example to scatter data to different processors from the one with rank 0

5.3.2.1 Scatter Python Objects

The example program executing the scatter is showcased next

```

1  #!/usr/bin/env python
2  from mpi4py import MPI
3
4  # Communicator
5  comm = MPI.COMM_WORLD
6
7  # Number of processes in the communicator group
8  size = comm.Get_size()
9
10 # Get the rank of the current process in the communicator group
11 rank = comm.Get_rank()
12
13 # Process with rank 0 gets a list with the data to be scattered
14 if rank == 0:
15     data = [(i + 1) ** 2 for i in range(size)]
16 else:
17     data = None
18
19 # Print data in each process before scattering
20 print(f'before scattering, data on rank {rank} is: {data}')
21
22 # Scattering occurs
23 data = comm.scatter(data, root=0)

```

```
24
25 # Print data in each process after scattering
26 print(f'after scattering, data on rank {rank} is: {data}')
```

Executing `mpirun -n 4 python scatter.py` yields:

```
1 before scattering, data on rank 2 is None
2 before scattering, data on rank 3 is None
3 before scattering, data on rank 0 is [1, 4, 9, 16]
4 before scattering, data on rank 1 is None
5 data for rank 2 is 9
6 data for rank 1 is 4
7 data for rank 3 is 16
8 data for rank 0 is 1
```

The members of the list from process 0 have been successfully scattered among the rest of the processes in the communicator group.

5.3.2.2 Scatter from Python Memory In the following example, we scatter a NumPy array among the processes in the communicator group by using the uppercase version of the method `comm.Scatter()`.

```
1 #!/usr/bin/env python
2 import numpy as np
3 from mpi4py import MPI
4
5 # Communicator
6 comm = MPI.COMM_WORLD
7
8 # Number of processes in the communicator group
9 size = comm.Get_size()
10
11 # Get the rank of the current process in the communicator group
12 rank = comm.Get_rank()
13
14 # Data to be sent
15 sendbuf = None
16
17 # Process with rank 0 populates sendbuf with a 2-D array,
18 # based on the number of processes in our communicator group
19 if rank == 0:
20     sendbuf = np.zeros([size, 10], dtype='i')
21     sendbuf.T[:, :] = range(size)
22
23     # Print the content of sendbuf before scattering
24     print(f'sendbuf in 0: {sendbuf}')
25
26 # Each process gets a buffer (initially containing just zeros)
27 # to store scattered data.
28 recvbuf = np.zeros(10, dtype='i')
29
```

```

30 # Print the content of recvbuf in each process before scattering
31 print(f'recvbuf in {rank}: {recvbuf}')
32
33 # Scattering occurs
34 comm.Scatter(sendbuf, recvbuf, root=0)
35
36 # Print the content of sendbuf in each process after scattering
37 print(f'Buffer in process {rank} contains: {recvbuf}')
```

Executing `mpiexec -n 4 python npscatter.py` yields:

```

1 recvbuf in 1: [0 0 0 0 0 0 0 0 0 0]
2 recvbuf in 2: [0 0 0 0 0 0 0 0 0 0]
3 recvbuf in 3: [0 0 0 0 0 0 0 0 0 0]
4 sendbuf in 0: [[0 0 0 0 0 0 0 0 0 0]
5                [1 1 1 1 1 1 1 1 1 1]
6                [2 2 2 2 2 2 2 2 2 2]
7                [3 3 3 3 3 3 3 3 3 3]]
8 recvbuf in 0: [0 0 0 0 0 0 0 0 0 0]
9 Buffer in process 2 contains: [2 2 2 2 2 2 2 2 2 2]
10 Buffer in process 0 contains: [0 0 0 0 0 0 0 0 0 0]
11 Buffer in process 3 contains: [3 3 3 3 3 3 3 3 3 3]
12 Buffer in process 1 contains: [1 1 1 1 1 1 1 1 1 1]
```

As we can see, the values in the 2-D array at process with rank 0, have been scattered among all our processes in the communicator group, based on their rank value.

5.3.3 Gather

The gather function is the inverse function to scatter. Data from each process is gathered in consecutive order based on the rank of the processor.

5.3.3.1 Gather Python Objects In this example, data from each process in the communicator group is gathered in the process with rank 0.




Figure 4: Example to gather data to different processors from the one with rank 0


```
1 #!/usr/bin/env python
2 from mpi4py import MPI
3
4 # Communicator
5 comm = MPI.COMM_WORLD
6
7 # Number of processes in the communicator group
8 size = comm.Get_size()
9
10 # Get the rank of the current process in the communicator group
11 rank = comm.Get_rank()
12
13 # Each process gets different data, depending on its rank number
14 data = (rank + 1) ** 2
15
16 # Print data in each process
17 print(f'before gathering, data on rank {rank} is: {data}')
18
19 # Gathering occurs
20 data = comm.gather(data, root=0)
21
22 # Process 0 prints out the gathered data, rest of the processes
23 # print their data as well
24 if rank == 0:
25     print(f'after gathering, process 0's data is: {data}')
26 else:
27     print(f'after gathering, data in rank {rank} is: {data}')
```

Executing `mpiexec -n 4 python gather.py` yields:

```
1 before gathering, data on rank 2 is 9
2 before gathering, data on rank 3 is 16
3 before gathering, data on rank 0 is 1
4 before gathering, data on rank 1 is 4
5 after gathering, data in rank 2 is None
6 after gathering, data in rank 1 is None
7 after gathering, data in rank 3 is None
8 after gathering, process 0's data is [1, 4, 9, 16]
```

The data from processes with rank 1 to `size - 1` have been successfully gathered in process 0.

5.3.3.2 Gather from Python Memory The example showcases the use of the uppercase method `comm.Gather()`. NumPy arrays from the processes in the communicator group are gathered into a 2-D array in process with rank 0.

```
1 #!/usr/bin/env python
2 import numpy as np
3 from mpi4py import MPI
4
5 # Communicator group
```

```

6 comm = MPI.COMM_WORLD
7
8 # Number of processes in the communicator group
9 size = comm.Get_size()
10
11 # Get the rank of the current process in the communicator group
12 rank = comm.Get_rank()
13
14 # Each process gets an array with data based on its rank.
15 sendbuf = np.zeros(10, dtype='i') + rank
16
17 # Print the data in sendbuf before gathering
18 print(f'Buffer in process {rank} before gathering: {sendbuf}')
19
20 # Variable to store gathered data
21 recvbuf = None
22
23 # Process with rank 0 initializes recvbuf to a 2-D array containing
24 # only zeros. The size of the array is determined by the number of
25 # processes in the communicator group
26 if rank == 0:
27     recvbuf = np.zeros([size, 10], dtype='i')
28
29     # Print recvbuf
30     print(f'recvbuf in process 0 before gathering: {recvbuf}')
31
32 # Gathering occurs
33 comm.Gather(sendbuf, recvbuf, root=0)
34
35 # Print recvbuf in process with rank 0 after gathering
36 if rank == 0:
37     print(f'recvbuf in process 0 after gathering: \n{recvbuf}')

```

Executing `mpiexec -n 4 python npgather.py` yields:

```

1 Buffer in process 2 before gathering: [2 2 2 2 2 2 2 2 2 2]
2 Buffer in process 3 before gathering: [3 3 3 3 3 3 3 3 3 3]
3 Buffer in process 0 before gathering: [0 0 0 0 0 0 0 0 0 0]
4 Buffer in process 1 before gathering: [1 1 1 1 1 1 1 1 1 1]
5 recvbuf in process 0 before gathering:
6 [[0 0 0 0 0 0 0 0 0 0]
7  [0 0 0 0 0 0 0 0 0 0]
8  [0 0 0 0 0 0 0 0 0 0]
9  [0 0 0 0 0 0 0 0 0 0]]
10 recvbuf in process 0 after gathering:
11 [[0 0 0 0 0 0 0 0 0 0]
12  [1 1 1 1 1 1 1 1 1 1]
13  [2 2 2 2 2 2 2 2 2 2]
14  [3 3 3 3 3 3 3 3 3 3]]

```

The values contained in the buffers from the different processes in the group have been gathered in the 2-D array in process with rank 0.

5.3.4 Allgather Memory Objects

This method is a many-to-many communication operation, where data from all processors is gathered in a continuous memory object on each of the processors. This is functionally equivalent to

1. A gather on rank 0
2. A Scatter from rank 0

However, this operation naturally has a performance bottleneck while all communication goes through rank0. Instead, we can use parallel communication between all of the processes at once to improve the performance. The optimization is implicit, and the user does not need to worry about it.

We demonstrate its use in the following example. Each process in the communicator group computes and stores values in a NumPy array (row). For each process, these values correspond to the multiples of the process' rank and the integers in the range of the communicator group's size. After values have been computed in each process, the different arrays are gathered into a 2D array (table) and distributed to ALL the members of the communicator group (as opposed to a single member, which is the case when `comm.Gather()` is used instead).




Figure 5: Example to gather the data from each process into ALL of the processes in the group

```

1 #!/usr/bin/env python
2 import numpy as np
3 from mpi4py import MPI
4
5 # Communicator group
6 comm = MPI.COMM_WORLD
7
8 # Number of processes in the communicator group
9 size = comm.Get_size()
10
11 # Get the rank of the current process in the communicator group
12 rank = comm.Get_rank()
13
14 # Initialize array and table
15 row = np.zeros(size)

```

```
16 table = np.zeros((size, size))
17
18 # Each process computes the local values and fills its array
19 for i in range(size):
20     j = i * rank
21     row[i] = j
22
23 # Print array in each process
24 print(f'Process {rank} table before Allgather: {table}\n')
25
26 # Gathering occurs
27 comm.Allgather([row, MPI.INT], [table, MPI.INT])
28
29 # Print table in each process after gathering
30 print(f'Process {rank} table after Allgather: {table}\n')
```

Executing

```
1 $ mpiexec -n 4 python allgather_buffer.py
```

results in the output similar to

```
1 Process 1 table before Allgather: [[0. 0.][0. 0.]]
2 Process 0 table before Allgather: [[0. 0.][0. 0.]]
3 Process 1 table after Allgather:  [[0. 0.][0. 1.]]
4 Process 0 table after Allgather:  [[0. 0.][0. 1.]]
```

As we see, after `comm.Allgather()` is called, every process gets a copy of the full multiplication table.

We have not provided an example for the Python object version as it is essentially the same and can easily be developed as an exercise.

5.4 Process Management

5.4.1 Dynamic Process Management with spawn

So far, we have focused on MPI used on a number of hosts that are automatically creating the process when `mpirun` is used. However, MPI also offers the ability to spawn a process in a communicator group. This can be achieved by using a spawn communicator and command.

Using

```
1 MPI.COMM_SELF.Spawn
```

will create a child process that can communicate with the parent. In the spawn code example, the manager broadcasts an array to the worker.

In this example, we have two Python programs: the first one being the manager and the second being the worker.




Figure 6: Example to spawn a program and start it on the different processors from the one with rank 0

```

1  #!/usr/bin/env python
2  from mpi4py import MPI
3  import numpy
4  import sys
5  import psutil
6
7  comm = MPI.COMM_SELF.Spawn(sys.executable,
8                             args=['mpi-worker.py'],
9                             maxprocs=(psutil.cpu_count(logical=False) -
10                                     1))
11
12 N = numpy.array(100, 'i')
13 comm.Bcast([N, MPI.INT], root=MPI.ROOT)
14 PI = numpy.array(0.0, 'd')
15 comm.Reduce(None, [PI, MPI.DOUBLE],
16             op=MPI.SUM, root=MPI.ROOT)
17 print(PI)
18 comm.Disconnect()

```

```

1  #!/usr/bin/env python
2  from mpi4py import MPI
3  import numpy
4
5  comm = MPI.Comm.Get_parent()
6  size = comm.Get_size()
7  rank = comm.Get_rank()
8
9  N = numpy.array(0, dtype='i')
10 comm.Bcast([N, MPI.INT], root=0)
11 h = 1.0 / N; s = 0.0
12 for i in range(rank, N, size):
13     x = h * (i + 0.5)
14     s += 4.0 / (1.0 + x**2)
15 PI = numpy.array(s * h, dtype='d')
16 comm.Reduce([PI, MPI.DOUBLE], None,

```

```
17         op=MPI.SUM, root=0)
18
19 comm.Disconnect()
```

To execute the example which calculates the number pi, please go to the examples directory and run the mpi-manager program with `-n 1`. There must only be 1 process because the additional processes are automatically created within the mpi-manager. The number of processes is automatically calculated according to the number of cores available minus 1 (because one core is already dedicated to the manager).

```
1 $ cd examples/spawn
2 $ mpiexec -n 1 python mpi-manager.py
```

This will result in an output close to the following:

```
1 3.1416009869231245
```

5.4.2 Futures

Futures is an mpi4py module that runs processes in parallel for intercommunication between such processes. The following Python program creates a visualization of a Julia set by utilizing the Futures modules, specifically via MPIPoolExecutor.

```
1 from mpi4py.futures import MPIPoolExecutor
2 import matplotlib.pyplot as plt
3 import numpy as np
4 from cloudmesh.common.StopWatch import StopWatch
5 from cloudmesh.common.variables import Variables
6 import multiprocessing
7
8 StopWatch.start("Overall time")
9
10 v = Variables()
11
12 if (v["multiplier"]):
13     multiplier = int((v["multiplier"]))
14     print(f"Proceeding since multiplier exists: {multiplier}")
15     pass
16 else:
17     print("No multiplier was input so multiplier defaults to 1\n"
18           "Use `cms set multiplier=2` to output higher resolution "
19           "Julia set image")
20     multiplier = 1
21     pass
22 if (v["workers"]):
23     workers = int((v["workers"]))
24     print(f"Proceeding since workers exists: {workers}")
25     pass
26 else:
```

```

27     print("No number of workers was input so workers defaults to 1\n"
28           "We suggest you use", multiprocessing.cpu_count(),
29           "workers for shortest runtime because that is the number of"
30           "threads you have available. Do this by issuing command "
31           f"$ cms set workers={multiprocessing.cpu_count()}`")
32     workers = 1
33     pass
34
35     x0, x1, w = -2.0, +2.0, 640*multiplier
36     y0, y1, h = -1.5, +1.5, 480*multiplier
37     dx = (x1 - x0) / w
38     dy = (y1 - y0) / h
39
40     c = complex(0, 0.65)
41
42
43     def julia(x, y):
44         z = complex(x, y)
45         n = 255
46         while abs(z) < 3 and n > 1:
47             z = z**2 + c
48             n -= 1
49         return n
50
51
52     def julia_line(k):
53         line = bytearray(w)
54         y = y1 - k * dy
55         for j in range(w):
56             x = x0 + j * dx
57             line[j] = julia(x, y)
58         return line
59
60
61     if __name__ == '__main__':
62         with MPIPoolExecutor(max_workers=workers) as executor:
63             image = executor.map(julia_line, range(h))
64             image = np.array([list(l) for l in image])
65             plt.imsave("julia.png", image)
66
67     Stopwatch.stop("Overall time")
68     Stopwatch.benchmark()

```

To run the program, issue this command in Git Bash:

```

1 $ cms set multiplier=2
2 $ cms set workers=4
3 $ mpiexec -n 1 python julia-futures.py

```

Note: if command cms is not found, make sure to install cloudmesh-common, cloudmesh_base, cloudmesh-inventory via pip

The multiplier variable serves as an integer which multiplies the standard resolution of the Julia set picture, which is 640x480. For example, issuing `cms set multiplier=3` will produce a 1920x1440 photo since 640x480 times 3 is 1920x1440. Not issuing a `cms set` command will cause the program to default to a multiplier of 1. The higher this number, the slower the runtime.

The workers variable serves as an integer which sets the number of workers to spawn for collaborative program execution. Not exporting this variable will cause it to default to 1 worker. The higher this number, the faster the runtime (up until the maximum number of threads on the CPU is surpassed).

The futures feature only works with `mpiexec -n 1` because it uses a method similar to that of spawn. Any other number will only repeat the program needlessly; it will not run faster or more efficiently.

The program will output a png image of a Julia set upon successful execution.

We created the numba version of this program in an attempt to achieve faster runtimes. Numba utilizes the jit decorator. For further explanation of numba, please see the Monte Carlo section of this document.

```
1 from mpi4py.futures import MPIPoolExecutor
2 import matplotlib.pyplot as plt
3 import numpy as np
4 from numba import jit
5 from cloudmesh.common.StopWatch import StopWatch
6 from cloudmesh.common.variables import Variables
7 import multiprocessing
8
9 StopWatch.start("Overall time")
10
11 v = Variables()
12
13 if (v["multiplier"]):
14     multiplier = int((v["multiplier"]))
15     print(f"Proceeding since multiplier exists: {multiplier}")
16     pass
17 else:
18     print("No multiplier was input so multiplier defaults to 1\n"
19           "Use `cms set multiplier=2` to output higher resolution "
20           "Julia set image")
21     multiplier = 1
22     pass
23 if (v["workers"]):
24     workers = int((v["workers"]))
25     print(f"Proceeding since workers exists: {workers}")
26     pass
27 else:
28     print("No number of workers was input so workers defaults to 1\n"
29           "We suggest you use", multiprocessing.cpu_count(),
30           "workers for shortest runtime because that is the number of "
31           "threads you have available. Do this by issuing command "
32           f"`cms set workers={multiprocessing.cpu_count()}`")
33     workers = 1
```



```

34     pass
35
36     x0, x1, w = -2.0, +2.0, 640*multiplier
37     y0, y1, h = -1.5, +1.5, 480*multiplier
38     dx = (x1 - x0) / w
39     dy = (y1 - y0) / h
40
41     c = complex(0, 0.65)
42
43
44     @jit(nopython=True)
45     def julia(x, y):
46         z = complex(x, y)
47         n = 255
48         while abs(z) < 3 and n > 1:
49             z = z**2 + c
50             n -= 1
51         return n
52
53
54     def julia_line(k):
55         line = bytearray(w)
56         y = y1 - k * dy
57         for j in range(w):
58             x = x0 + j * dx
59             line[j] = julia(x, y)
60         return line
61
62
63     if __name__ == '__main__':
64         with MPIPoolExecutor(max_workers=workers) as executor:
65             image = executor.map(julia_line, range(h))
66             image = np.array([list(l) for l in image])
67             plt.imsave("julia.png", image)
68
69     Stopwatch.stop("Overall time")
70     Stopwatch.benchmark()

```

No Jit	1 Worker	2 Workers	6 Workers	12 Workers	20 Workers
640x480	22.470 s	12.220 s	4.946 s	4.384 s	5.257 s
1280x960	45.951 s	23.702 s	8.982 s	6.258 s	6.523 s
1920x1440	68.779 s	34.652 s	12.933 s	8.385 s	8.042 s

Jit Enabled	1 Worker	2 Workers	6 Workers	12 Workers	20 Workers
640x480	24.551 s	12.499 s	5.632 s	5.054 s	6.616 s

Jit Enabled	1 Worker	2 Workers	6 Workers	12 Workers	20 Workers
1280x960	46.183 s	23.406 s	9.543 s	7.190 s	8.226 s
1920x1440	68.366 s	34.569 s	12.938 s	9.278 s	8.854 s

- These benchmark times were generated using a Ryzen 5 3600 CPU with 16 GB RAM on a Windows 10 computer. The Ryzen 5 3600 is a 6-core, 12-thread processor.

No Jit	1 Worker	2 Workers	3 Workers	4 Workers
640x480	51.555 s	49.103 s	48.501 s	48.983 s
1280x960	66.044 s	56.652 s	53.693 s	52.929 s
1920x1440	87.918 s	68.069 s	61.836 s	59.414 s

- These benchmark times were generated using a Raspberry Pi 4 Model B 2018 with 8 GB RAM on a Raspbian 10 (codename buster) OS. It uses an ARM Cortex-A72 4-core, 4-thread processor. On this Raspberry Pi, 4 workers can be used via the `cms set workers=4` and `mpirun -np 1 --oversubscribe python julia-futures.py` commands. However, any number of workers greater than 4 causes the program to hang and timeout with an unknown MPI spawn error, likely because the Pi does not support using a greater number of threads. Also, numba cannot be used on Pi.

Jit does not appear to shorten the program runtimes, causing it to be longer in most instances except for a few higher resolution outputs.

6 MPI Example Programs

In this section, we will showcase to you some simple MPI example programs.

6.1 MPI Ring Example

The MPI Ring example program is one of the classical programs every MPI programmer has seen. Here a message is sent from the manager to the workers while the processors are arranged in a ring, and the last worker sends the message back to the manager. Instead of just doing this once, our program does it multiple times and adds every time a communication is done 1 to the integer send around. Figure 1 showcases the process graph of this application.




Figure 7: Processes organized in a ring perform a sum operation

In the example, the user provides an integer that is transmitted from the process with rank 0, to process with rank 1, and so on until the data returns to process 0. Each process increments the integer by 1 before transmitting it to the next one, so the final value received by process 0 after the ring is complete is the sum of the original integer plus the number of processes in the communicator group.

```

1  #!/usr/bin/env python
2  # USAGE: mpiexec -n 4 python ring.py --count 1000
3  from mpi4py import MPI
4  import click
5  from cloudmesh.common.StopWatch import StopWatch
6
7  @click.command()
8  @click.option('--count', default=1, help='Number of messages send.')
9  @click.option('--debug', default=False, help='Set debug.')
10 def ring(count=1, debug=False):
11     comm = MPI.COMM_WORLD # Communicator
12     rank = comm.Get_rank() # Get the rank of the current process
13     size = comm.Get_size() # Get the size of the communicator group
14     if rank == 0:
15         print(f'Communicator group with {size} processes')
16         data = int(input('Enter an integer to transmit: ')) # Input
17         # the data
18         data += 1 # Data is
19         # modified
20     if rank == 0: # ONLY processor 0 uses the stopwatch
21         StopWatch.start(f"ring {size} {count}")
22     for i in range(0, count):
23         if rank == 0:
24             comm.send(data, dest=rank + 1) # send data to neighbor
25             data = comm.recv(data, source=size - 1)
26             if debug:
27                 print(f'Final data received in process 0: {data}')
28         elif rank == size - 1:
29             data = comm.recv(source=rank - 1) # recieve data from
30             # neighbor
31             data += 1 # Data is modified
32             comm.send(data, dest=0) # Sent to process 0,
33             # closing the ring
34         elif 0 < rank < size - 1:
35             data = comm.recv(source=rank - 1) # recieve data from
36             # neighbor
37             data += 1 # Data is modified

```

```
33         comm.send(data, dest=rank + 1)      # send to neighbor
34     if rank == 0:
35         print(f'Final data received in process 0: {data}')
36         assert data == count * size          # verify
37     if rank == 0:
38         Stopwatch.stop(f"ring {size} {count}") #print the time
39         Stopwatch.benchmark()
40
41 if __name__ == '__main__':
42     ring()
```

Executing the code in the example by entering `mpirun -n 2 python ring.py` in the terminal will produce the following result:

```
1 Communicator group with 4 processes
2 Enter an integer to transmit: 6
3 Process 0 transmitted value 7 to process 1
4 Process 1 transmitted value 8 to process 2
5 Process 2 transmitted value 9 to process 3
6 Process 3 transmitted value 10 to process 0
7 Final data received in process 0 after ring is completed: 10
```

As we can see, the integer provided to process 0 (6 in this case) was successively incremented by each process in the communicator group to return a final value of 10 at the end of the ring.

6.2 Counting Numbers

The following program generates arrays of random numbers each 20 (n) in length with the highest number possible being 10 (max_number). It then uses a function called `count()` to count the number of 8's in each data set. The number of 8's in each list is stored `count_data`. `Count_data` is then summed and printed out as the total number of 8's.

The program allows you to set the program parameters. Note that the program has on purpose a bug in it as it does not communicate the values `m`, `max_number`, or `find` with a broadcast from rank 0 to all workers. Your task is to modify and complete this program.

```
1 # Run with
2 #     mpirun -n 4 python count.py
3
4 # To change the values set them on your terminal on the
5 # machine running rank 0 with
6
7 # export N=20
8 # export MAX=10
9 # export FIND=8
10
11 # Assignment:
12 # Add to this code the broadcast of the 3 parameters to all workers
13
```

```

14 import os
15 import random
16 from mpi4py import MPI
17
18 # Get the input values or set them to a default
19 n = int(os.environ.get("N") or 20)
20 max_number = int(os.environ.get("MAX") or 10)
21 find = int(os.environ.get("FIND") or 8)
22
23
24 comm = MPI.COMM_WORLD # Communicator
25 size = comm.Get_size() # Number of processes
26 rank = comm.Get_rank() # Rank of this process
27
28 # Each process gets different data, depending on its rank number
29 data = []
30 for i in range(n):
31     r = random.randint(1, max_number)
32     data.append(r)
33 count = data.count(find)
34
35 print(rank, count, data) # Print data from each process
36 count_data = comm.gather(count, root=0) # Gather the data
37
38 # Process 0 prints out the gathered data, rest of the processes
39 if rank == 0:
40     print(rank, count_data)
41     total = sum(count_data)
42     print(f"Total number of {find}'s:", total)

```

Executing `mpiexec -n 4 python count.py` gives us:

```

1 1 1 [7, 5, 2, 1, 5, 5, 5, 4, 5, 2, 6, 5, 2, 1, 8, 7, 10, 9, 5, 6]
2 3 3 [9, 2, 9, 8, 2, 7, 7, 2, 10, 1, 2, 5, 3, 5, 10, 8, 10, 10, 8, 10]
3 2 3 [1, 3, 8, 5, 7, 8, 4, 2, 8, 5, 10, 7, 10, 1, 6, 5, 9, 6, 6, 7]
4 0 3 [6, 9, 10, 2, 4, 8, 8, 9, 4, 1, 6, 8, 6, 9, 7, 5, 5, 6, 3, 4]
5
6 0 [3, 1, 3, 3]
7
8 Total number of 8's: 10

```

6.3 Monte Carlo Calculation of Pi

A very nice example to showcase the potential for doing lots of parallel calculations is to calculate the number pi. This is quite easily achieved while using a Monte Carlo Method.

We start with the mathematical formulation of the Monte Carlo calculation of pi. For each quadrant of the unit square, the area is pi. Therefore, the ratio of the area outside of the circle is pi over four. With this in mind, we can use the Monte Carlo Method for the calculation of pi.

The following is a visualization of the program's methodology to calculate pi:




Figure 8: montecarlographic

The following `montecarlo.py` program generates an estimation of pi using the methodology and equation shown above. Increasing the total number of iterations will increase the accuracy.

```

1 import random as r
2 import math as m
3 import time
4
5 start = time.time()
6
7 inside = 0                                # Number of darts that land inside.
8 trials = 100000                           # Number of Trials.
9
10 for i in range(0, trials):                # Iterate for the number of darts.
11     x2 = r.random()**2                     # Generate random x, y in [0, 1]
12     y2 = r.random()**2
13
14     if m.sqrt(x2 + y2) < 1.0:              # Increment if inside unit circle.
15         inside += 1
16
17 # inside / trials = pi / 4
18 pi = (float(inside) / trials) * 4
19 end = time.time()
20
21 print(pi)                                # Value of pi found
22 print(end - start)                       # Execution time

```

Instead of running this on one processor, we can run the calculation on many. Implicitly this increases the accuracy while running more trials at the same time as we run them all in parallel. Overhead does exist by starting the MPI program and gathering the result. However, if the trial number is large enough, it is negligible.

The following program shows the MPI implementation [7]:

```

1 # Originally from https://cvw.cac.cornell.edu/python/exercise
2 # Modified by the cloudmesh team

```

```
3 """
4 An estimate of the numerical value of pi via Monte Carlo integration.
5 Computation is distributed across processors via MPI.
6 """
7
8 import numpy as np
9 from mpi4py import MPI
10 import matplotlib
11 matplotlib.use('Agg')
12 import matplotlib.pyplot as plt
13 import sys
14 from cloudmesh.common.StopWatch import StopWatch
15
16 StopWatch.start("Overall time")
17 def throw_darts(n):
18     """
19     returns an array of n uniformly random (x,y) pairs lying within the
20     square that circumscribes the unit circle centered at the origin,
21     i.e., the square with corners at (-1,-1), (-1,1), (1,1), (1,-1)
22     """
23     darts = 2*np.random.random((n,2)) - 1
24     return darts
25
26 def in_unit_circle(p):
27     """
28     returns a boolean array, whose elements are True if the
29     corresponding
30     point in the array p is within the unit circle centered at the
31     origin,
32     and False otherwise -- hint: use np.linalg.norm to find the length
33     of a vector
34     """
35     return np.linalg.norm(p,axis=-1)<=1.0
36
37 def estimate_pi(n, block=100000):
38     """
39     returns an estimate of pi by drawing n random numbers in the square
40     [[-1,1], [-1,1]] and calculating what fraction land within the unit
41     circle;
42     in this version, draw random numbers in blocks of the specified
43     size,
44     and keep a running total of the number of points within the unit
45     circle;
46     by throwing darts in blocks, we are spared from having to allocate
47     very large arrays (and perhaps running out of memory), but still
48     can get
49     good performance by processing large arrays of random numbers
50     """
51     total_number = 0
52     i = 0
53     while i < n:
54         if n-i < block:
55             block = n-i
```

```

49         darts = throw_darts(block)
50         number_in_circle = np.sum(in_unit_circle(darts))
51         total_number += number_in_circle
52         i += block
53     return (4.*total_number)/n
54
55 def estimate_pi_in_parallel(comm, N):
56     """
57     on each of the available processes,
58     calculate an estimate of pi by drawing N random numbers;
59     the manager process will assemble all of the estimates
60     produced by all workers, and compute the mean and
61     standard deviation across the independent runs
62     """
63
64     if rank == 0:
65         data = [N for i in range(size)]
66     else:
67         data = None
68     data = comm.scatter(data, root=0)
69     #
70     pi_est = estimate_pi(N)
71     #
72     pi_estimates = comm.gather(pi_est, root=0)
73     if rank == 0:
74         return pi_estimates
75
76
77 def estimate_pi_statistics(comm, Ndarts, Nruns_per_worker):
78     results = []
79     for i in range(Nruns_per_worker):
80         result = estimate_pi_in_parallel(comm, Ndarts)
81         if rank == 0:
82             results.append(result)
83     if rank == 0:
84         pi_est_mean = np.mean(results)
85         pi_est_std = np.std(results)
86         return pi_est_mean, pi_est_std
87
88 if __name__ == '__main__':
89     """
90     for N from 4**5 to 4**14 (integer powers of 4),
91     compute mean and standard deviation of estimates of pi
92     by throwing N darts multiple times (Nruns_total times,
93     distributed across workers)
94     """
95     comm = MPI.COMM_WORLD
96     rank = comm.Get_rank()
97     size = comm.Get_size()
98     if rank == 0:
99         print("MPI size = {}".format(size))
100         sys.stdout.flush()
101     Nruns_total = 64

```



```

102     Nruns_per_worker = Nruns_total // size
103     #
104     estimates = []
105     for log4N in range(5,15):
106         N = int(4**log4N)
107         result = estimate_pi_statistics(comm, N, Nruns_per_worker)
108         if rank == 0:
109             pi_est_mean, pi_est_std = result
110             estimates.append((N, pi_est_mean, pi_est_std))
111             print(N, pi_est_mean, pi_est_std)
112             sys.stdout.flush()
113     if rank == 0:
114         estimates = np.array(estimates)
115         plt.figure()
116         plt.errorbar(np.log2(estimates[:,0]), estimates[:,1], yerr=
            estimates[:,2])
117         plt.ylabel('estimate of pi')
118         plt.xlabel('log2(number of darts N)')
119         plt.savefig('pi_vs_log2_N.png')
120         plt.figure()
121         plt.ylabel('log2(standard deviation)')
122         plt.xlabel('log2(number of darts N)')
123         plt.plot(np.log2(estimates[:,0]), np.log2(estimates[:,2]))
124         plt.savefig('log2_std_vs_log2_N.png')
125     MPI.Finalize()
126
127     StopWatch.stop("Overall time")
128     StopWatch.benchmark()

```

To run this program using git bash, change directory to the folder containing this program and issue the command:

```
1 $ mpiexec -n 4 python parallel_pi.py
```

The number after `-n` can be changed to however many cores one has on their processor.

However, running this program takes upwards of 4 minutes to complete with 6 cores. We can use numba to speed up the program execution time.

Additionally, we can run this program on multiple hosts. For instance, you can use a machinefile or rankfile to execute the program on a PI cluster.

6.3.1 Numba

Numba, an open-source JIT (just in time) compiler, is a Python module that translates Python code into machine code for faster runtimes.

The numba version of the Monte Carlo program runs faster, even cutting runtime down by a few minutes:

```
1 # Originally from https://cvw.cac.cornell.edu/python/exercise
2 # Modified by the cloudmesh team
3 from __future__ import print_function, division
4 """
5 An estimate of the numerical value of pi via Monte Carlo integration.
6 Computation is distributed across processors via MPI.
7 """
8
9 import numpy as np
10 from mpi4py import MPI
11 import matplotlib
12 matplotlib.use('Agg')
13 import matplotlib.pyplot as plt
14 import sys
15 from numba import jit
16 from cloudmesh.common.StopWatch import StopWatch
17
18 StopWatch.start("Overall time")
19 @jit(nopython=True)
20 def throw_darts(n):
21     """
22     returns an array of n uniformly random (x,y) pairs lying within the
23     square that circumscribes the unit circle centered at the origin,
24     i.e., the square with corners at (-1,-1), (-1,1), (1,1), (1,-1)
25     """
26     darts = 2*np.random.random((n,2)) - 1
27     return darts
28
29 def in_unit_circle(p):
30     """
31     returns a boolean array, whose elements are True if the
32     corresponding
33     point in the array p is within the unit circle centered at the
34     origin,
35     and False otherwise -- hint: use np.linalg.norm to find the length
36     of a vector
37     """
38     return np.linalg.norm(p,axis=-1)<=1.0
39
40 def estimate_pi(n, block=100000):
41     """
42     returns an estimate of pi by drawing n random numbers in the square
43     [[-1,1], [-1,1]] and calculating what fraction land within the unit
44     circle;
45     in this version, draw random numbers in blocks of the specified
46     size,
47     and keep a running total of the number of points within the unit
48     circle;
49     by throwing darts in blocks, we are spared from having to allocate
50     very large arrays (and perhaps running out of memory), but still
51     can get
52     good performance by processing large arrays of random numbers
53     """
```

```

46     """
47     total_number = 0
48     i = 0
49     while i < n:
50         if n-i < block:
51             block = n-i
52             darts = throw_darts(block)
53             number_in_circle = np.sum(in_unit_circle(darts))
54             total_number += number_in_circle
55             i += block
56     return (4.*total_number)/n
57
58 def estimate_pi_in_parallel(comm, N):
59     """
60     on each of the available processes,
61     calculate an estimate of pi by drawing N random numbers;
62     the manager process will assemble all of the estimates
63     produced by all workers, and compute the mean and
64     standard deviation across the independent runs
65     """
66
67     if rank == 0:
68         data = [N for i in range(size)]
69     else:
70         data = None
71     data = comm.scatter(data, root=0)
72     #
73     pi_est = estimate_pi(N)
74     #
75     pi_estimates = comm.gather(pi_est, root=0)
76     if rank == 0:
77         return pi_estimates
78
79 def estimate_pi_statistics(comm, Ndarts, Nruns_per_worker):
80     results = []
81     for i in range(Nruns_per_worker):
82         result = estimate_pi_in_parallel(comm, Ndarts)
83         if rank == 0:
84             results.append(result)
85     if rank == 0:
86         pi_est_mean = np.mean(results)
87         pi_est_std = np.std(results)
88         return pi_est_mean, pi_est_std
89
90 if __name__ == '__main__':
91     """
92     for N from 4**5 to 4**14 (integer powers of 4),
93     compute mean and standard deviation of estimates of pi
94     by throwing N darts multiple times (Nruns_total times,
95     distributed across workers)
96     """
97     comm = MPI.COMM_WORLD
98     rank = comm.Get_rank()

```

```

99     size = comm.Get_size()
100     if rank == 0:
101         print("MPI size = {}".format(size))
102         sys.stdout.flush()
103     Nruns_total = 64
104     Nruns_per_worker = Nruns_total // size
105     #
106     estimates = []
107     for log4N in range(5,15):
108         N = int(4**log4N)
109         result = estimate_pi_statistics(comm, N, Nruns_per_worker)
110         if rank == 0:
111             pi_est_mean, pi_est_std = result
112             estimates.append((N, pi_est_mean, pi_est_std))
113             print(N, pi_est_mean, pi_est_std)
114             sys.stdout.flush()
115     if rank == 0:
116         estimates = np.array(estimates)
117         plt.figure()
118         plt.errorbar(np.log2(estimates[:,0]), estimates[:,1], yerr=
119                     estimates[:,2])
120         plt.ylabel('estimate of pi')
121         plt.xlabel('log2(number of darts N)')
122         plt.savefig('pi_vs_log2_N.png')
123         plt.figure()
124         plt.ylabel('log2(standard deviation)')
125         plt.xlabel('log2(number of darts N)')
126         plt.plot(np.log2(estimates[:,0]), np.log2(estimates[:,2]))
127         plt.savefig('log2_std_vs_log2_N.png')
128     MPI.Finalize()
129     StopWatch.stop("Overall time")
130     StopWatch.benchmark()

```

Note how before the definition of functions in this code, there is the `@jit(nopython=True)` decorator, which translates each defined function into faster machine code. To install and use numba, simply issue the command `pip install numba` within a terminal. Here is the command to execute the numba version of the Monte Carlo program:

```
1 $ mpiexec -n 4 python parallel_pi_numba.py
```

Cores	parallel_pi.py execution time	parallel_pi_numba.py execution time
6	237.873 s	169.678 s
5	257.720 s	199.572 s
4	326.811 s	239.160 s
3	383.343 s	289.433 s
2	545.500 s	403.289 s

Cores	parallel_pi.py execution time	parallel_pi_numba.py execution time
1	1075.68 s	810.525 s

- These benchmark times were generated using a Ryzen 5 3600 CPU with 16 GB RAM on a Windows 10 computer.

Note: Please be advised that we use Cloudmesh.StopWatch which is a convenient program to measure time and display the details for the computer. However, it is not threadsafe and, at this time, only measures times in the second range. If your calculations for other programs are faster or the trial number is too slow, you can use other benchmarking methods.

6.3.2 Running Monte Carlo on multiple hosts

Another way to increase the performance of our program would be executing it on multiple hosts.

As an example, we can run the program in a cluster of 7 Pis: a manager PI4, and six worker PI3s. Be advised, however, that we do not use numba on RaspberryOS, hence execution can take a relatively long time in comparison to the numba version. For a reference, simultaneously running a copy of the program on the cluster (7 processes total) took around 40 minutes. However, using a machinefile to run four copies of the program on each node (for a total of 28 processes) significantly sped up execution, taking only a fourth of that time (~ 10 minutes).

First, we need to make sure that mpi4py is installed on all Pis. For that purpose, you can follow our tutorial to Deploy MPI for Python (mpi4py) on your Pi Cluster using Cloudmesh.

Next, we send a copy of the program to each of the hosts in our cluster. It is important that the file be stored in the same directory address for every host. For this example, we send it to the home directory ~/.

```

1 (ENV3) pi@red:~ $ for h in red red0{1..6}; do
2 > scp parallel_pi.py pi@$h:~/ &
3 > done
4 [4] 2176
5 [5] 2177
6 [6] 2178
7 [7] 2179
8 [8] 2180
9 [9] 2181
10 [10] 2182
11 (ENV3) pi@red:~ $

```

Next, we create a machinefile to specify the number of hosts and cores to be employed by mpi4py during execution. Notice that we are employing the four available cores on each node.

```

1 pi@red slots=4
2 pi@red01 slots=4
3 pi@red02 slots=4
4 pi@red03 slots=4
5 pi@red04 slots=4
6 pi@red05 slots=4
7 pi@red06 slots=4

```

We will need to save the machinefile only in the node from which the `mpiexec` command is executed.

Finally, we run the program by calling `mpiexec` from the command line. Note we have added the parameter `-machinefile` to specify the machinefile location. Additionally we used the full address of the Python binary from ENV3 to ensure that every host runs the program inside our environment.

```

1 (ENV3) pi@red:~ $ mpiexec -n 28 -machinefile ./machinefile ~/ENV3/bin/
  python parallel_pi.py
2 MPI size = 28
3 1024 3.1455775669642856 0.05692609648889646
4 4096 3.138096400669643 0.02867811562631138
5 16384 3.1421116420200894 0.01228796385009885
6 65536 3.1414925711495534 0.005532607075034393
7 262144 3.1418974740164622 0.0029191407443025902
8 1048576 3.141772815159389 0.0017866555629716318
9 4194304 3.1415172815322876 0.0008114422251076501
10 16777216 3.1416021159717014 0.0003635004457387496
11 67108864 3.1415649854711125 0.0001929914184235647
12 268435456 3.1416016641472067 9.638285643379315e-05
13 ...

```

6.4 Mandelbrot

We can run a program which outputs a visualization of a Mandelbrot data set, which, like the Julia set, is a fractal (the image repeats itself upon zooming in). This program runs processes in parallel and also has numba JIT decorators to achieve faster runtimes:

```

1 from matplotlib import pyplot
2 from mpi4py import MPI
3 import numpy
4 from numba import jit
5 from cloudmesh.common.StopWatch import StopWatch
6
7
8 @jit(nopython=True)
9 def mandelbrot(x, y, maxit):
10     c = x + y * 1j
11     z = 0 + 0j
12     it = 0
13     while abs(z) < 2 and it < maxit:
14         z = z ** 2 + c

```

```
15         it += 1
16     return it
17
18
19     x1, x2 = -2.0, 1.0
20     y1, y2 = -1.0, 1.0
21     w, h = 1200, 800
22     maxit = 127
23
24     comm = MPI.COMM_WORLD
25     size = comm.Get_size()
26     rank = comm.Get_rank()
27     if rank == 0:
28         Stopwatch.start(f'parallel {size}')
29     # number of rows to compute here
30     N = h // size + (h % size > rank)
31
32     # first row to compute here
33     start = comm.scan(N) - N
34
35     # array to store local result
36     Cl = numpy.zeros([N, w], dtype='i')
37
38     dx = (x2 - x1) / w
39     dy = (y2 - y1) / h
40     for i in range(N):
41         y = y1 + (i + start) * dy
42         for j in range(w):
43             x = x1 + j * dx
44             Cl[i, j] = mandelbrot(x, y, maxit)
45
46     # gather results at root (process 0)
47
48     counts = comm.gather(N, root=0)
49     C = None
50     if rank == 0:
51         C = numpy.zeros([h, w], dtype='i')
52
53     rowtype = MPI.INT.Create_contiguous(w)
54     rowtype.Commit()
55
56     comm.Gatherv(sendbuf=[Cl, MPI.INT],
57                  recvbuf=[C, (counts, None), rowtype],
58                  root=0)
59
60     rowtype.Free()
61
62     if rank == 0:
63         Stopwatch.stop(f'parallel {size}')
64
65         pyplot.imshow('mandelbrot-parallel-numba.png', C)
66         pyplot.imshow('mandelbrot-parallel-numba.pdf', C)
67
```

```

68     Stopwatch.benchmark()
69     # pyplot.imshow(C, aspect='equal')
70     # pyplot.show()

```

Like other programs, mandelbrot can be executed via `mpiexec -n 4 python mandelbrot-parallel-numba.py`, with the appropriate `-n` parameter according to the user's system.

At rank 0, the program starts and ends a benchmark for analysis of which `-n` parameter will give the shortest runtime.

		mandelbrot-parallel-numba.py execution
Cores	mandelbrot-parallel.py execution time	time
6	3.071 s	0.422 s
5	3.791 s	0.434 s
4	3.920 s	0.427 s
3	5.769 s	0.473 s
2	5.010 s	0.520 s
1	9.891 s	1.765 s

- These benchmark times were generated using a Ryzen 5 3600 CPU with 16 GB RAM on a Windows 10 computer.

This program will save an image and pdf called mandelbrot:




Figure 9: mandelbrot

6.4.1 Assignments

1. Use numba to speed up the code. Create a tutorial including installation instructions.
2. Display results with matplotlib as created by the picture
3. Modify cloudmesh.Stopwatch so we can use it for smaller time measurements

6.5 Other MPI Example Programs

You will find lots of example programs on the internet when you search for it. Please let us know about such examples and we will add the here. You can also contribute to our repository and add example programs that we then include in this document. In return you will become a co-author or get acknowledged.

- A program to calculate k-means is provided at
 - <https://medium.com/@hyeamykim/parallel-k-means-from-scratch-2b297466fdcd>

6.6 GPU Programming with MPI

In case you have access to computers with GPUs, you can naturally utilize them accordingly from Python with the appropriate GPU drivers.

In case not all have a GPU, you can use rankfiles to control the access and introduce through conditional programming based on rank access to the GPUs.

7 Parameter Management

Although this next topic is not directly related to MPI and mpi4py, it is very useful in general. Often we ask ourselves the question, “how do we pass parameters to a program, including MPI?” There are multiple ways to achieve this, for example, with environment variables, command-line arguments, and configuration files. We will explain each of these methods and provide simple examples.

7.1 Using the Shell Variables to Pass Parameters

`os.environ` in Python allows us to easily access environment variables that are defined in a shell. It returns a dictionary having the user’s environmental variable as key and their values as value.

To demonstrate its use, we have written a `count.py` program that uses `os.environ` to optionally pass parameters to an MPI program.

This example is included in a previous section named `Counting Numbers` and we like you to look it over.

If the user changed the value of N, MAX, or FIND in the terminal using, for example, `export FIND="5"` (shown below) `os.environ.get("FIND")` would set the find variable equal to 5.

```
1 $ export FIND="5"
2 $ mpiexec -n 4 python count.py
```

```

3 1 1 [6, 3, 3, 8, 4, 1, 1, 4, 4, 3, 8, 5, 10, 8, 8, 7, 2, 4, 1, 9]
4 3 0 [3, 1, 4, 1, 6, 4, 9, 3, 1, 8, 8, 6, 4, 3, 7, 1, 8, 6, 1, 1]
5 2 3 [5, 5, 4, 6, 8, 5, 9, 3, 7, 7, 10, 6, 7, 3, 2, 8, 3, 10, 7, 10]
6 0 3 [7, 8, 6, 9, 6, 7, 5, 6, 1, 2, 1, 2, 9, 5, 9, 8, 5, 1, 8, 1]
7 0 [3, 1, 3, 0]
8 Total number of 5's: 7

```

However, if the user does not define any environment variables, FIND will default to 8.

```

1 $ mpiexec -n 4 python count.py
2 1 0 [5, 5, 2, 6, 6, 3, 5, 3, 3, 2, 3, 9, 7, 1, 3, 7, 1, 7, 1, 3]
3 3 1 [7, 1, 5, 1, 2, 2, 10, 7, 2, 1, 2, 6, 4, 6, 10, 10, 5, 8, 10, 10]
4 2 0 [5, 1, 4, 4, 9, 9, 5, 1, 1, 3, 9, 3, 5, 2, 5, 7, 9, 7, 10, 5]
5 0 1 [6, 6, 5, 6, 4, 10, 3, 5, 5, 2, 5, 2, 7, 6, 7, 8, 5, 7, 6, 4]
6 0 [1, 0, 0, 1]
7 Total number of 8's: 2

```

Assignment:

1. One thing we did not do is use the broadcast method to properly communicate the 3 environment variables. We like you to improve the code and submit to us.

Let us assume we use the Python program

```

1 from cloudmesh.common.StopWatch import StopWatch
2 from time import sleep
3 import os
4
5 n=int(os.environ["N"])
6 StopWatch.start(f"processors {n}")
7 sleep(0.1*n)
8 print(n)
9 StopWatch.stop(f"processors {n}")
10 StopWatch.benchmark()

```

This Python program does not set a variable N on its own. It refers to os.environ which is a module that refers to variables exported within the same shell that executes the program.

Setting the variable/parameter can either be done via the export shell command such as

```
1 $ export N=8
```

or while passing the parameter in the same line as a command such as demonstrated next

```
1 $ N=1; python environment-parameter.py
```

This can be generalized while using a file with many different parameters and commands. For example, placing this in a file called `run.sh` with these contents:

```

1 $ N=1; python environment-parameter.py
2 $ N=2; python environment-parameter.py

```

It allows us to execute the programs sequentially in the file with

```
1 $ sh run.py
```

In our case, we are also using cloudmesh.StopWatch to allow us easily to fgrep for the results we may be interested in to conduct benchmarks. Here is an example workflow to achieve this

```
1 # This command creates an environment variable called N
2 $ export N=10
3 # This command prints the environment variable called N
4 $ echo $N
5 # This command launches a Python environment
6 $ python -i
7 >>> import os
8 >>> os.environ["N"]
9 >>> exit()
10 $ python environment-parameter.py
11 $ sh run.sh
12 $ sh run.sh | fgrep "csv,processors"
```

7.1.1 Using click to pass parameters

Click is a convenient mechanism to define parameters that can be passed via options to python programs. To showcase its use please inspect the following program

```
1 import click
2 from cloudmesh.common.StopWatch import StopWatch
3 from time import sleep
4 import os
5
6 @click.command()
7 @click.option('--n', default=1, help='Number of processors.')
8 def work(n):
9     n=int(n)
10    StopWatch.start(f"processors {n}")
11    sleep(0.1*n)
12    print(n)
13    StopWatch.stop(f"processors {n}")
14    StopWatch.benchmark()
15
16 if __name__ == '__main__':
17     work()
```

You can manually set the variable in git bash in the same line as you open the .py file

```
1 $ python click-parameter.py --n=3
```

8 SLURM

In case you run long running jobs it is often useful to have access to a batch queuing system. Such a batch queue enables one to submit the jobs to a queue and they are scheduled for execution based on a scheduling policy. One such framework is SLURM. We describe how to use mpi4py from a batch queuing system with SLURM. Slurm stands for **S**imple **L**inux **U**tility for **R**esource **M**anagement. It is an open-source job scheduler for a compute cluster to carry out tasks efficiently and in a particular order while using the cluster's resources. SLURM supports batch jobs but also allows the use of resources in interactive mode. SLURM is a popular batch queueing system used on many advanced supercomputers. However, it is possible to install and use SLURM on a cluster of Raspberry Pis. SLURM can utilize mpi4py to achieve a unique processing power that replaces the use of individual computer threads with entire computers in and of themselves.

8.1 Installation of SLURM on a Raspberry Pi Cluster

The installation takes around an hour on a cluster of four Raspberry Pi 4 Model B computers.

To use the cloudmesh SLURM command, one must have cloudmesh installed by using the following commands.

We assume you are in a venv Python environment. Ours is called (ENV3)

```
1 (ENV3) you@yourlaptop $ mkdir ~/cm
2 (ENV3) you@yourlaptop $ cd ~/cm
3 (ENV3) you@yourlaptop $ pip install cloudmesh-installer
4 (ENV3) you@yourlaptop $ cloudmesh-installer get pi
```

Initialize the cms command:

```
1 (ENV3) you@yourlaptop $ cms help
```

Then clone the cloudmesh-slurm repository:

```
1 (ENV3) you@yourlaptop $ cd ~/cm
2 (ENV3) you@yourlaptop $ cloudmesh-installer get cmd5
3 (ENV3) you@yourlaptop $ git clone https://github.com/cloudmesh/
  cloudmesh-slurm.git
4 (ENV3) you@yourlaptop $ cd cloudmesh-slurm
5 (ENV3) you@yourlaptop $ pip install -e .
6 (ENV3) you@yourlaptop $ cms help
```

You may proceed if `slurm` shows in the documented commands.

After following the burn tutorial and ensuring that the cluster is online, you have two methods of installing SLURM.

8.1.1 Method 1 - Install from Host

You can install SLURM on a cluster by executing commands from the host computer. The host computer is the same computer that is previously burned your SD Cards and is referred to as `you@yourlaptop`. This machine can be used to `ssh` into each of the Pis.

To install it, use the command:

```
1 cms slurm pi install as host --hosts=red,red0[1-3] --mount=//dev//sda
```

The `mount` parameter is meant to have double slashes no matter the OS of the host.

The `--hosts` parameter needs to include the hostnames of your cluster, including manager and workers, separated by comma using a parameterized naming scheme.

The `--mount` parameter points to the mount place of your USB, inserted in the top-most blue USB3.0 port (on Pi 4's) on your manager PI.

WARNING: This USB drive **will be formatted** and all data on it will be erased.

The command will take a long time to finish. It may appear to not progress at certain points, but please be patient. However they will last hopefully not longer than 45 minutes. The reason this takes such a long time is that at time of writing of this tutorial, the prebuilt SLURM packages did not work, so we compile it from source.

Once the script completes, you can check if SLURM is installed by issuing on the manager:

```
srun --nodes=3 hostname
```

and replacing the `--nodes` parameter with the number of workers.

You will see an output similar to

```
1 (ENV3) pi@red:~ $ srun --nodes=3 hostname
2 red01
3 red02
4 red03
5 (ENV3) pi@red:~ $
```

The nodes may be out of order. That is okay and normal.

8.1.2 Method 2 - Install on Manager

The manager Pi is the designated Raspberry Pi computer that will act as the central headquarters of the entire cluster. The manager runs the `slurmd` daemon, which is the controller of all the nodes and their jobs. In our documentation, our example manager is named `red`.

This method is for those who do not want to use a host computer to facilitate the installation; instead, the installation is run directly on the manager Pi. However, this method is more tedious as the user must reconnect to the Pi after it reboots to rerun the script (three times in total).

8.1.2.1 Install cloudmesh on Manager Pi This method involves the user logging into the manager via `ssh` and first installing cloudmesh in the manager with:

```
1 (ENV3) you@yourhostcomputer $ ssh red
2 pi@red $ curl -Ls http://cloudmesh.github.io/get/pi | sh -
```

This output is printed upon successful installation.

```
1 Please activate with
2
3     source ~/ENV3/bin/activate
4
5 Followed by a reboot
```

After activating venv with the source command and rebooting via `sudo reboot`, issue the commands:

```
1 (ENV3) you@yourhostcomputer $ ssh red
2 pi@red:~ $ cd ~/cm
3 pi@red:~/cm $ git clone https://github.com/cloudmesh/cloudmesh-slurm.
  git
4 pi@red:~/cm $ cd cloudmesh-slurm
5 pi@red:~/cm/cloudmesh-slurm $ pip install -e .
6 pi@red:~/cm/cloudmesh-slurm $ cms help
```

The slurm command should appear in the list.

8.1.2.2 Install SLURM on Manager Pi Run this command to begin SLURM installation:

```
1 pi@red:~/cm/cloudmesh-slurm $ cms slurm pi install --workers=red0[1-3]
  --mount=/dev/sda
```

The user must `ssh` back into the manager after the cluster reboots and perform the last command (`cms slurm pi install...`) 3 more times. The script will inform the user when this is no longer necessary and SLURM is fully installed.

Notice this method does not need two forward slashes in `--mount` because it is done on Raspberry Pi OS and not Windows. It can only be done on Raspberry Pi OS because the method is purposefully done on the manager Pi, to begin with.

You can check if SLURM is installed by issuing on the manager:

```
srun --nodes=3 hostname
```

and replacing the `--nodes` parameter with the number of workers.

You will see an output similar to

```
1 (ENV3) pi@red:~ $ srun --nodes=3 hostname
2 red01
3 red02
4 red03
5 (ENV3) pi@red:~ $
```

The nodes may be out of order. That is okay and normal.

8.2 Install SLURM on a Single Raspberry Pi

Instead of installing SLURM on an entire cluster, let us now consider the case in which you only have one Raspberry Pi. To make job management simple on this Pi, we can install SLURM on that one computer. This one computer has no workers and is a manager to its own self. The user can make and automate jobs for simplicity's sake, and the same computer will carry out those jobs.

Single-node installation, which is a SLURM cluster with only one node, can be easily configured by using the `host` command with the manager and workers listed as the same hostname. In the following example, `red` is the single-node.

```
1 cms slurm pi install as host --hosts=red,red --mount=//dev//sda
```

8.3 MPI Example

To run a test MPI example, `ssh` into the manager and then use the `example` command. This is only possible if `cms` is installed on the Pi; if you have not done this because you installed SLURM via the `host` method, then refer to section 3.1 to install cloudmesh on Pi. Then run the following (change the number after `--n` to the number of nodes):

```
1 (ENV3) you@yourhostcomputer $ ssh red
2 pi@red:~ $ cms slurm pi example --n=3
```

This `cms slurm` command runs `salloc -N 3 mpiexec python -m mpi4py.bench helloworld` but the number after `-N` is altered to whatever is input for the `--n` parameter. Do not run the `salloc` command. It is unnecessary when we have already implemented it within the aforementioned `cms slurm pi example` command. It is just listed here for reference. The output will be similar to:

```
1 pi@red:~ $ cms slurm pi example --n=3
2 salloc: Granted job allocation 17
3 Hello, World! I am process 0 of 3 on red01.
4 Hello, World! I am process 1 of 3 on red02.
```

```
5 Hello, World! I am process 2 of 3 on red03.  
6 salloc: Relinquishing job allocation 17
```

9 Links to Other Documents

Here are a couple of links that may be useful. We have not yet looked over them but include them.

- <https://research.computing.yale.edu/sites/default/files/files/mpi4py.pdf>
- <https://www.nesi.org.nz/sites/default/files/mpi-in-python.pdf>
- <https://www.kth.se/blogs/pdc/2019/08/parallel-programming-in-python-mpi4py-part-1/>
- <http://education.molssi.org/parallel-programming/03-distributed-examples-mpi4py/index.html>
- <http://www.cec-hpc.be/assets/training/mpi4py.pdf>
- <https://www.csc.fi/documents/200270/224366/mpi4py.pdf/825c582a-9d6d-4d18-a4ad-6cb6c43fef8>

9.1 Assignment

1. Review the resources and provide a short summary that we add to this document above the appropriate link

10 Appendix

10.1 Git Bash on Windows

Git bash is a implementation of the bash shell for windows that also includes Git.

Git is an open-source software which helps to manage repository version control, particularly with GitHub repos.

To verify whether you have Git in the first place, you can press **Win + R** on your desktop, type **cmd**, and press **Enter**. Then type **git clone** and press **Enter**. If you do not have Git installed, it will say **'git' is not recognized as an internal or external command...**

As long as Git does not change the structure of their website and hyperlinks, you should be able to download Git from here and skip to Step #2: <https://git-scm.com/downloads>

1. Open a web browser and search **git**. Look for the result that is from **git-scm.com** and click Downloads.

2. Click **Download for Windows**. The download will commence. Open the file once it is finished downloading.
3. The UAC Prompt will appear. Click **Yes** because Git is a safe program. It will show you Git's license: a GNU General Public License. After understanding the terms, click **Next**. 1. The GNU General Public License means that the program is open-source (free of charge).
4. Click **Next** to confirm that `C:\Program Files\Git` is the directory where you want Git to be installed.
5. Click **Next** unless you would like an icon for Git on the desktop (in which case you can check the box and then click **Next**).
6. Click **Next** to accept the text editor, click **Next** again to Let Git decide the default branch name, click **Next** again to run Git from the command line and 3rd party software, click **Next** again to use the OpenSSL library, click **Next** again to checkout Windows-style, click **Next** again to use MinTTY, click **Next** again to use the default git pull, click **Next** again to use the Git Credential Manager Core, click **Next** again to enable file system caching, and then click **Install** because the experimental features are not necessary.
7. Wait for the green progress bar to finish. Congratulations— you have installed Git and Git Bash. You can now run it as an administrator by pressing the Windows key, typing `git bash`, right clicking **Git Bash**, and clicking **Run as administrator**. Click **Yes** in the UAC prompt that appears.

10.2 Make on Windows

Makefiles provide a good feature to organize workflows while assembling programs or documents to create an integrated document. Within `makefiles` you can define targets that you can call and are then executed. Preconditions can be used to execute rules conditionally. This mechanism can easily be used to define complex workflows that require a multitude of interdependent actions to be performed. Makefiles are executed by the program `make` that is available on all platforms.

On Linux, it is likely to be pre-installed, while on macOS you can install it with Xcode. On Windows, you have to install it explicitly. We recommend that you install `gitbash` first. After you install `gitbash`, you can install `make` from an administrative `gitbash` terminal window. To start one, go to the search field next to the Windows icon on the bottom left and type in `gitbash` without a **RETURN**. You will then see a selection window that includes **Run as administrator**. Click on it. As you run it as administrator, it will allow you to install `make`. The following instructions will provide you with a guide to install `make` under windows.

10.2.1 Installation

Please visit

- <https://sourceforge.net/projects/ezwinports/files/>

and download the file

- 'make-4.3-without-guile-w32-bin.zip'

After the download, you have to extract and unzip the file as follows in a gitbash that you started as an administrative user:




Figure 10: administrativegitbash

Figure: Screenshot of opening gitbash in admin shell

```
1 $ cp make-4.3-without-guile-w32-bin.zip /usr
2 $ cd /usr
3 $ unzip make-4.3-without-guile-w32-bin.zip
```

Now start a new terminal (a regular non-administrative one) and type the command

```
1 $ which make
```

It will provide you the location if the installation was successful

```
1 /usr/bin/make
```

to make sure it is properly installed and in the correct directory.

10.3 Installing WSL on Windows 10

WSL is a layer that allows the running of Linux executables on a Windows machine.

To install WSL2 your computer must have Hyper-V support enabled. This does not work on Windows Home, and you need to upgrade to Windows Pro, Edu, or some other Windows 10 version that supports it. Windows Edu is typically free for educational institutions. The Hyper-V must be enabled from your BIOS, and you need to change your settings if it is not enabled.

More information about WSL is provided at

- <https://docs.microsoft.com/en-us/windows/wsl/install-win10>

To install WSL2, you can follow these steps while using Powershell as an administrative user and run

```
1 ps$ dism.exe /online /enable-feature /featurename:Microsoft-Hyper-V-All /all /norestart
2 ps$ dism.exe /online /enable-feature /featurename:Microsoft-Windows-Subsystem-Linux /all /norestart
3 ps$ dism.exe /online /enable-feature /featurename:VirtualMachinePlatform /all /norestart
4 ps$ wsl --set-default-version 2
```

The next command will restart your computer so make sure that all your files and applications are saved:

```
1 ps$ Restart-Computer
```

Windows will say that it is working on updates (enabling the features). Once logging back in, download this msi file, open it and complete the installation to update WSL:

- https://wslstorestorage.blob.core.windows.net/wslblob/wsl_update_x64.msi

Once the installation is complete, download and install the Ubuntu 20.04 LTS image from the Microsoft store:

- <https://www.microsoft.com/en-us/p/ubuntu/9nblggh4msv6?activetab=pivot:overviewtab>

and click Launch.

Run Ubuntu and create a username and passphrase.

Make sure not just to give an empty passphrase but choose a secure one.

Next run in a new instance of elevated (admin) Powershell:

```
1 ps$ wsl.exe --set-version Ubuntu 2
```

Now you can use the Ubuntu distro freely. The WSL2 application will be in your shortcut menu in [Start](#). You can launch this WSL2 and install MPI on it by referring to the Ubuntu installation instructions at the beginning of this document. The same number of cores and threads will be available to use in the [mpiexec](#) command as the number of cores and threads on the host computer.

10.4 Benchmarks

This section is in more detail published at this link. If the link does not work use this Link.

We explain how we can manage long-running benchmarks. There are many useful tools to conducting benchmarks such as `timeit`, `cprofile`, `line_profiler`, and `memry_profiler` to name only a few. However, we present here an extremely easy way to obtain runtimes while dealing with the fact that they could run multiple hours or even days and could cause your system to crash. Hence if we wor to run it in a single program it will lead to a loss of information and many hours of unneeded replication.

We use and demonstrate how we achieve this with a simple StopWatch, creation of shell scripts and even the integration of Jupyter notebooks.

10.4.1 Prerequisites

As usual, we recommend that you use a virtual env. dependent on where your python 3 is installed, please adapt accordingly (python, or python3). Also, test out which version of python you have. On Windows, we assume you have gitbash installed and use it.

```
1 $ python3 -- version # observe that you have the right version
2 $ python3 -m venv ~/ENV
3 $ source ~/ENV3/bin/activate
4 # or for Windows gitbash
5 # source ~/ENV3/Scripts/activate
```

10.4.2 System Parameters

It is essential that we benchmark programs to show their effect on the time consumed to obtain the results. Various factors play a role. This includes the number of physical computers involved, the number of processors on each computer, the number of cores on each computer, and the number of threads for each core. We can summarise these parameters as a vector such as

```
1 S(N, p, c, t)
```

Where

- `S` = is a placeholder for the system
- `N` = Number of computers or nodes
- `p` = Number of processors per node
- `c` = Number of cores per processor
- `t` = Number of threads per processor

In some cases, it may be more convenient to specify the total values as

```
1 S^T(N, N*p, N*p*c, N*p*c*t)
```

and

- T = indicates total

In the case of heterogeneous systems, we define multiple such vectors to form a list of vectors.

For the rest of the section, we assume the system is homogeneous.

10.4.2.1 System Information Cloudmesh provides an easy command that can be used to obtain information to derive these values while using the command. However, it only works if the number of processors on the same node is 1.

```
1 pip install cloudmesh-cmd5
2 cms help      # call it after the install as it sets some defaults
3 cms sysinfo
```

The output will be looking something like

Attribute	Value
cpu	Intel(R) Core(TM) i7-7920HQ CPU @ 3.10GHz
cpu_cores	4
cpu_count	8
cpu_threads	8
frequency	scpufreq(current=3100, min=3100, max=3100)
mem.active	5.7 GiB
mem.available	5.8 GiB
mem.free	96.7 MiB
mem.inactive	5.6 GiB
mem.percent	63.7 %
mem.total	16.0 GiB
mem.used	8.2 GiB
mem.wired	2.4 GiB
platform.version	10.16
python	3.9.5 (v3.9.5:0a7dcdbd13, ...)
python.pip	21.1.2
python.version	3.9.5
sys.platform	darwin
uname.machine	x86_64
uname.node	mycomputer
uname.processor	i386
uname.release	20.5.0
uname.system	Darwin
uname.version	Darwin Kernel Version 20.5.0:
user	gregor

To obtain the vectors you can say

```
1 cms sysinfo -v
2 cms sysinfo -t
```

where `-v` specifies the vector and `-t` the totals. Knowing these values will help you structure your benchmarks.

10.4.2.2 Parameters A benchmark is typically run while iterating over a number of parameters and measuring some system parameters that are relevant for the benchmark, such as the runtime of the program or application.

Let us assume our application is called `f` and its parameters are `x` and `y`

To create benchmarks over `x` and `y` we can generate them in various ways.

10.4.2.3 Python only solution For all programs, we will store the output of the benchmarks in a directory called `benchmark`. Please create it.

```
1 $ mkdir benchmark
```

you may be able to run your benchmark simply as a loop this is especially the case for smaller benchmarks.

```
1 import pickle
2 from cloudmesh.common.StopWatch import StopWatch
3
4 def f(x,y, print_benchmark=False, checkpoint=True):
5     # run your application with values x and y
6     print (f"Calculate f({x},{y})")
7     StopWatch.start(f"f{x},{y}")
8     result = x*y
9     StopWatch.stop(f"f{x},{y}")
10    if print_benchmark:
11        StopWatch.benchmark()
12    if checkpoint:
13        pickle.dump(result, open(f"benchmark/f-{x}-{y}.pkl", "wb" ))
14    return result
15
16 x_min = 0
17 x_max = 2
18 d_x = 1
19 y_min = 0
20 y_max = 1
21 d_y = 1
22 for x in range(x_min, x_max, dx):
23     for y in range(y_min, y_max, dy):
24         # run the function with parameters
25         result = f(x ,y, print_benchmark=True)
```

10.4.2.4 Script solution In some cases, the functions themselves may be large and in case the benchmark causes a crash of the python program executing it we would have to start over. In such

cases, it is better to develop scripts that take parameters so we can execute the program through shell scripts and exclude those that fail.

For this, we rewrite the python program via command-line arguments that we pass along.

```

1 # stored in file f.py
2 import click
3
4 @click.command()
5 @click.option('--x', default=20, help='The x value')
6 @click.option('--y', default=40, help='The y value')
7 @click.option('--print_benchmark', default=True, help='prints the
    benchmark result')
8 @click.option('--checkpoint', default=True, help='Creates a checkpoint'
    )
9 f(x,y, print_benchmark=False, checkpoint=True):
10     ... see previous program
11     return result
12
13 if __name__ == '__main__':
14     f()

```

Now we can run this program with

```
1 $ python f.py --x 10 --y 5
```

To generate now the different runs from the loop we can do it either via Makefiles or write a program creating commands where we produce a script listing each invocation. Let us call this program `sweep-generator.py`.

```

1 x_min = 0
2 x_max = 2
3 d_x = 1
4 y_min = 0
5 y_max = 1
6 d_y = 1
7 for x in range(x_min, x_max, dx):
8     for y in range(y_min, y_max, dy):
9         print (f"cms banner f({x}, {y}); "
10              f"python f.py --x {x} --y {y}")

```

The result will be

```

1 cms banner f(0,0); python f.py --x 0 --y 0
2 ...

```

and so on. The banner will print a nice banner before you execute the real function so it is easier to monitor when execution

To create a shell script, simply redirect it into a file such as

```
1 $ python sweep-generator.py sweep.sh
```

Now you can execute it with

```
1 $ sh sweep.sh | tee result.log
```

The `tee` command will redirect the output to the file `result`, while still reporting its progress on the terminal. In case you want to run it without monitoring or `tee` is not supported properly you just run it as

```
1 $ sh sweep.sh >result.log
```

In case you need to monitor the progress for the latter you can use

```
1 $ tail -f result.log
```

The advantage of this approach is that you can in case of a failure identify which benchmarks succeeded and exclude them from your next run of `sweep.sh` so you do not have to redo them. This may be useful if you identify that you ran out of resources for a parameterized run and it crashed.

10.4.2.5 Integrating timers The beauty about cloudmesh is that it has built-in timers and if properly used we can use them even across different invocations of the function `f`.

we simply have to `fgrep` to the log file to extract the information in the `csv` lines with

```
1 fgrep "#csv" result.log
```

This can then be further post-processed.

Cloudmesh also includes a `cloudmesh.Shell.cm_grep`, `cloudmesh.common.readfile`, and other useful functions to make the processing of shell scripts and their output easier.

10.4.2.6 Integration of Jupyter Notebooks Jupyter notebooks provide a simple mechanism to prototype. However, how do we now integrate them into a benchmarking suite? Certainly, we can just create the loop in the notebooks conducting the parameter sweep, but in case of a crash, this becomes highly unscalable.

So what we have to do is augment a notebook so that we can

1. pass along the parameters,
2. execute it from the command line.

For this, we use `papermill` that allows us to just do these two tasks. Install it with

```
1 pip install papermill
```


Then when you open up jupyter-lablab and import our code. Create a new cell. In this cell you place all parameters for your run that you like to modify such as

```
1 x = 0
2 y = 0
```

This cell can be augmented with a tag called “parameters”. To do this open the “cog” and enter in the tag name “parameters”. Make sure you save the tag and the notebook. Now we can use `papermill` to run our notebook with parameters such as

```
1 $ mkdir benchmark
2 $ papermill sweep.ipynb benchmark/sweep-0-0.ipynb --x 0 --y 0 | tee
   benchmark/result-0-0.log
3 ...
```

Naturally, we can auto-generate this as follows

```
1 x_min = 0
2 x_max = 2
3 d_x = 1
4 y_min = 0
5 y_max = 1
6 d_y = 1
7 for x in range(x_min, x_max, dx):
8     for y in range(y_min, y_max, dy):
9         print (f"cms banner f({x}, {y}); "
10              f"papermill sweep.ipynb benchmark/sweep-{x}-{y}.ipynb"
11              f"      --x {x} --y {y}"
12              f"      | tee benchmark/result-{x}-{y}.log")
```

This will produce a series of commands that we can also redirect into a shell script and then execute

10.4.3 Combining the logs

As we have the logs all in the benchmark directory, we can even combine them and select the `csv` lines with

```
1 $ cat benchmark/*.log | fgrep "#csv"
```

Now you can apply further processing such as importing it into pandas or any other spreadsheet-like tools you like to use for the analysis.

11 Assignments

1. Develop a section explaining what MPI-IO is
2. Develop a section to explain Collective I/O with NumPy arrays.

3. Add a section on how to use Numpy with MPI, including the installation of NumPy. This is not to have a tutorial about numpy, but how to use numpy within mpi4py. Subtasks include
 1. Download Numpy with `pip install numpy` in a terminal
 2. `import numpy as np` to use NumPy in the program
 3. Explain the advantages of NumPy over pickled lists
 - Numpy stores memory contiguously
 - Uses a smaller number of bytes
 - Can multiply arrays by index
 - It is faster
 - Can store different data types, including images
 - Contains random number generators
 4. Add a specific, very small tutorial on using some basic numpy features as they may be useful for MPI application development. This may include the following and be added to the appendix
 1. To define an array type: `np.nameofarray([1,2,3])`
 2. To get the dimension of the array: `nameofarray.ndim`
 3. To get the shape of the array (the number of rows and columns): `nameofarray.shape`
 4. To get the type of the array: `nameofarray.dtype`
 5. To get the number of bytes: `nameofarray.itemsize`
 6. To get the number of elements in the array: `nameofarray.size`
 7. To get the total size: `nameofarray.size * nameofarray.itemsize`

Please, note that we have a very comprehensive tutorial on NumPy and there is no point to repeat that, we may just point to it and improve that tutorial where needed instead.

4. Convert the parallel rank program from <https://mpitutorial.com/tutorials/performing-parallel-rank-with-mpi/> to mpi4py. Write a tutorial for it.
5. Develop tutorials that showcase multiple communicators and groups. See <https://mpitutorial.com/tutorials/introduction-to-groups-and-communicators/>
6. Complete the count example while adding a broadcast to it to communicate the parameters. Provide a modified tutorial.
7. Test out the machinefile, host, and rankfile section. Improve if needed.

References

- [1] G. von Laszewski, “Python for Cloud Computing,” Indiana University, Bloomington IN, U.S.A., Online Book, Feb. 2020 [Online]. Available: <https://cloudmesh-community.github.io/pub/vonLaszewski-python.pdf>
- [2] “MPICH: High-Performance Portable MPI.” Sep-2021 [Online]. Available: <https://www.mpich.org>
- [3] “Open MPI: Open Source High Performance Computing.” Sep-2021 [Online]. Available: <https://www.open-mpi.org>
- [4] “Intel MPI Library,” *Intel*. Sep-2021 [Online]. Available: <https://software.intel.com/content/www/us/en/develop/tools/oneapi/components/mpi-library.html#gs.c5q095>
- [5] “Microsoft MPI - Message Passing Interface.” Sep-2021 [Online]. Available: <https://docs.microsoft.com/en-us/message-passing-interface/microsoft-mpi>
- [6] “MPI Solutions for GPUs,” *NVIDIA Developer*. May-2021 [Online]. Available: <https://developer.nvidia.com/mpi-solutions-gpus>
- [7] “Cornell Virtual Workshop: Exercise: Monte Carlo with mpi4py.” Sep-2021 [Online]. Available: <https://cvw.cac.cornell.edu/python/exercise>